


Trinity S³AI



TERNARY NETWORK FLOATS

A datapath with no multiplier anywhere: two formats close the two halves of a ternary node, one sign-select primitive builds both a radio and a transformer at zero DSP, and the multiply-free path is exact

Dmitrii Vasilev

ORCID 0009-0008-4294-6159

admin@t27.ai

11 August 2026

Abstract

A ternary node has three sites that could carry a number format; one needs to. The weight is a code: multiplying by it is a sign-select. The sample arrives ADC-native. The accumulator alone has a range to spend, and Ternary Network Floats spend it — a sign, a balanced-ternary exponent of E_t trits, an M -bit mantissa, $1 + E_t + M = N$.

The weight is **GFternary**, a two-bit $\{-\varphi, 0, +\varphi\}$ alphabet, and it is where the multiplier is decided. Removing one rather than cheapening it requires the weight lattice to be *closed* under weight application and accumulation alike. $\mathbb{Z}[\varphi]$ is, because $\varphi^2 = \varphi + 1$: applying a weight is the Fibonacci step $(a, b) \mapsto (b, a + b)$, accumulation is componentwise, and a layer's linear path is exact. Machine-checked in Coq.

Closure does not single out φ ; enumeration does. A scale is multiplier-free exactly when it is an algebraic integer whose companion matrix has entries in $\{0, \pm 1\}$. Degree two admits φ alone, and nothing lies between the shift and φ . The family $r^d = r + 1$ reaches any granularity at *one addition* and d registers: fineness costs registers, not adders.

Among the twenty-one formats carrying their own range — those a training run can live in — GFternary ranks highest on throughput per area in this XC7A200T post-route comparison, at 0.1797 MHz/LUT against 0.1631 for the next entry and 0.0295 for **posit32**, with nine of the twenty-one ours. Formats needing an external learned scale are tabulated apart. On an XC7A200T the folded ternary neuron costs 28.0 LUT per weight at fan-in 8, rising to 33.1 and 33.2 at fan-in 16 and 32, with *no DSP at any fan-in*.

Which rung follows from the budget — φ at four bits, $r^5 = r^3 + 1$ at five, $r^6 = r + 1$ at six — reaching 3.7% of **fp32** at six bits and one adder. On the block axis the element is closed,

the best eight-level codebook beating a deployed one by 0.9%, and the scale is not: a φ grid at four bits beats MXFP4 on two models at 0.125 fewer bits per weight.

Five retractions are marked in place below.

1 What this paper argues

The results below were found in the order a measurement campaign finds things. They compose into one argument, and the argument is easier to follow stated first.

A ternary node has three sites where a number format could live, and only one of them needs one. The weight is a code and multiplying by it is a sign-select; the sample arrives ADC-native; the accumulator is the only object carrying a range that must be spent. So the design question is not which format to use but which *two* objects to close — an alphabet for the weight and a float for the accumulator — and Sections 2 to 6 settle the second by a law: once a range is named, the split between exponent and mantissa is determined rather than chosen, because the error depends on the mantissa width alone. Inverted, the same law reads a format’s behaviour back out of its round-trip error and resolves an 83-format catalogue into four staircase forms and no fifth.

TNF is one corner of a square, and Section 18 draws it. Two rules fix an exponent budget — the golden-ratio rule $E = \text{round}((N - 1)/\varphi^2)$ and the width rule $1 + E + M = N$ — and two radices encode it, binary or balanced ternary. The four families are GF, GF-T, BNF and TNF, each a distinct format rather than a renaming of another, and at eight bits the two rules turn out to name the identical layout.

The alphabet is where the multiplier is decided, and the deciding property is algebraic. Section 19.3 shows that removing a multiplier — as against making one cheaper — requires the weight lattice to be *closed* under the operation that applies a weight. $\mathbb{Z}[\varphi]$ is, because $\varphi^2 = \varphi + 1$, so a weight application is the Fibonacci step and an accumulation is componentwise integer addition; Fibonacci integers are not, because $F_4 F_4 = 9$ is not a Fibonacci number, and a datapath built on them must carry a stage returning products to the representable set. That stage measures an order of magnitude more area than the entire closed accumulation, and pipelining does not rescue it: a normalisation cascade cannot be cut below one compare-and-subtract per stage, so its floor lies above a closed path’s single addition. The closure half of the argument is machine-checked.

Closure does not single out φ by itself, and Section 19.4 says what it does single out. A scale is applicable without a multiplier exactly when it is an algebraic integer whose companion matrix has entries in $\{0, \pm 1\}$, which is a finite condition and therefore enumerable rather than arguable. Degree one gives the shift; degree two gives φ and nothing else; degree three gives three more, finer. Between the shift and φ there is nothing, and φ is the only multiply-free refinement of the power-of-two ladder available at two registers. The hierarchy continues — and stops being cheap at degree four, where the coefficient vector loses its sparsity and the cost doubles.

Which rung to use is then an empirical question with a law behind it (Section 20). A ladder of ratio r with n magnitudes spans r^{n-1} , so fine steps are bought with range, and the optimum moves down the hierarchy as the code budget grows: the shift at three bits, φ at four, the plastic number at five, on two models with the same ordering throughout. The optimum has a closed form in the weight histogram, transferring between models to four decimals, and the one place it errs — the four-bit pair — is corrected by weighting the error with the loss curvature rather than the weight magnitude.

Section 21 reports where none of this wins. On a short block with a shared scale the optimal eight-level codebook beats a deployed one by 0.9%, so the axis is not an element-format problem at all, and it belongs to MXFP4. We state it as the competition’s result and give the bound that makes a fourth attempt of ours unwarranted.

Two sections are about cost rather than kind. Section 13 makes area and precision commensurable, so a variable-length exponent field can be priced in mantissa bits of silicon; Section 14 then finds that for a logarithmic format the field layout is the smaller half of that cost, and that we had been measuring the wrong half.

Five of our own claims were falsified by the measurements meant to support them, and each is retracted where it was made rather than quietly dropped. Section 27 generalises the shape those errors took, because in almost every case the measurement was correct and the comparison around it was not.

2 The format

$$\text{TNF16} = [s \mid E=4 \text{ balanced-ternary trits} \mid M=11 \text{ bits}],$$

$$v = (-1)^s \left(1 + \frac{M}{29}\right) 2^e, \quad e = \sum_i t_i 3^i \in [-40, 40]. \quad (1)$$

Three properties follow from the layout rather than from tuning.

No regime decode. The dominant cost in a tapered format is the variable-length regime: it must be located, its length counted, and the significand barrel-shifted into place. That cost is paid on any fabric, ternary included. Fixed fields make it a bit slice.

The exponent is already ternary. Four balanced-ternary trits give $3^4 = 81$ exponent steps. On a ternary fabric the exponent add is native — no binary carry, no conversion between bases.

Precision does not taper. Nine mantissa bits at every magnitude, where a tapered format narrows towards the extremes. This is the mechanism behind the accuracy result of Section 5: the win appears where the comparison format has spent its mantissa on range.

3 Why a ladder, and why this one

A format that is used as a reference has different requirements from one used as a workhorse. A workhorse should be fast and accurate where the data is. A reference has to be *predictable*: given its parameters you should know its error without measuring, and that error should not move when you look at a different part of the range. A ruler whose divisions grow towards one end is not a ruler.

Table 1 tests both properties on this ladder.

The precision is a closed form, and it is derived rather than fitted. Write $u = 2^{-(M+1)}$ for the unit roundoff. Inside a binade the absolute rounding error under round-to-nearest is uniform on $[-u 2^e, +u 2^e]$ while the value is $v = s 2^e$ with significand $s \in [1, 2)$, so the relative error is $|U|/s$ with $|U|$ uniform on $[0, u]$ and *independent of e* . Hence

Theorem 1 (Precision law). *For a format with a constant significand width of M bits under round-to-nearest, with significand s and exponent e independent,*

$$\mathbb{E}[|rel\ err|] = \frac{1}{2} \mathbb{E}[1/s] \cdot 2^{-(M+1)},$$

independently of e .

Table 1: The precision law and the flatness, measured in exact rationals. “Flatness” is the ratio of the worst magnitude band to the best: 1.000 would mean the error does not depend on magnitude at all.

rung	M	$2^{-(M+1)}$	measured	ratio	flatness
TNF8	4	3.12e−2	1.22e−2	0.39	1.000
TNF16	9	9.77e−4	3.60e−4	0.37	1.070
TNF32	25	1.49e−8	5.49e−9	0.37	1.070
TNF64	52	1.11e−16	4.22e−17	0.38	1.050
TNF128	115	1.20e−35	4.44e−36	0.37	1.070
TNF256	242	7.07e−74	2.69e−74	0.38	1.050
TNF512	497	1.22e−150	4.51e−151	0.37	1.070
TNF1024	1006	7.29e−304	2.77e−304	0.38	1.050

The exponent drops out entirely — that is the flatness of the next paragraph, proved rather than observed. The constant is not universal: it is $\frac{1}{2} \ln 2 = 0.3466$ for a significand uniform on $[1, 2)$ and $\frac{1}{2} (2 \ln 2)^{-1} = 0.3607$ under a logarithmic (Benford) significand. Our workload has $\mathbb{E}[1/s] = 0.7721$, predicting 0.3861; the measured figures in Table 1 average 0.3756 over eight rungs, with a spread of 0.369–0.390 that brackets it.

The practical consequence is the one a reference needs: given M and the significand distribution of a workload, the error of a rung nobody has built is known before it is built.

The error does not move with magnitude, and cannot. The derivation above removes e from the expression, so magnitude-independence is a property of the layout rather than a happy measurement. Empirically flatness stays between 1.05 and 1.07 — at most a 7% difference between the closest and furthest parts of a workload spanning 76 binary decades. This is the property fixed fields buy and the one a tapered format cannot have: spending mantissa bits on range is exactly what makes precision depend on where you are.

What this does not claim. Not that TNF is the most accurate format at any given width — Section 6 shows posit16 and binary16 ahead near unity, and Section 6 shows the binary formats ahead outright at 32 bits. A reference does not need to win those comparisons. It needs to be the thing you measure them against.

4 The law is general, and it measures tapering

Section 3 derived the error of a fixed-field format. Nothing in that derivation is specific to TNF, so it should hold for any format whose significand width does not vary — and should visibly fail for one whose width does. That makes it an instrument rather than a description, and this section uses it as one.

Definition 1 (Effective mantissa). For a format and a magnitude band B , define

$$M_{\text{eff}}(B) = -\log_2 \left(\frac{2 \mathbb{E}[|\text{rel err}| \mid B]}{\mathbb{E}[1/s]} \right) - 1,$$

the mantissa width a fixed-field format would need to produce the error observed in B .

Theorem 2 (Diagnostic). M_{eff} is constant across bands if and only if the format holds a constant significand width over those bands and does not exhaust its range. For a tapered format M_{eff} decreases with $|e|$, and the slope $dM_{\text{eff}}/d|e|$ is the taper rate in bits of significand per binade.

Table 2: Effective mantissa by magnitude band, and the fitted taper rate. Declared M is recovered to 0.01 bits for every fixed-field format, which validates the law; the tapered formats separate cleanly.

format	declared M	$ e $ 0–6	6–12	12–20	20–30	bits/binade
TNF16	9	8.99	8.99	9.00	9.01	+0.001
GF16	9	8.99	8.99	9.00	9.01	+0.001
bfloat16	7	7.04	7.04	7.00	7.04	−0.000
binary16	10	10.01	10.01	7.41	0.52	—
posit16	—	10.72	9.34	7.48	5.17	− 0.254
takum16	—	9.62	8.17	7.34	7.04	− 0.113

Table 2 applies it. The measurement is a uniform significand on $[1, 2)$ so that $\mathbb{E}[1/s] = \ln 2$ exactly, removing the workload from the question.

The law recovers what the formats declare. TNF16 and GF16 return 8.99–9.01 against a declared 9; bfloat16 returns 7.00–7.04 against a declared 7. A law that reconstructs a format’s own parameter to a hundredth of a bit, from nothing but its round-trip error, is doing more than curve-fitting.

Tapering becomes a number. posit16 sheds 0.261 bits of significand per binade ($R^2 = 0.999$) and takum16 appears to shed 0.113 over the thirty binades of this sweep — but see Section 4.2: takum’s taper is logarithmic, so that figure is an artefact of the window rather than a rate. That is the quantity a tapered format trades for range, and to our knowledge it is not usually reported as a single figure. It also explains the accuracy tables directly — posit16 starts 1.7 bits ahead of TNF16 and ends 3.8 behind.

Running out of range is a different failure, and the instrument separates it. binary16 is a fixed-field format, yet its M_{eff} collapses from 10.01 to 0.52. It is not tapering: it is subnormalising and then overflowing. A constant slope indicates a taper; a cliff indicates a range limit. The two look alike in an accuracy table and do not look alike here.

Corollary 1. *A format’s position on the range–precision trade can be stated as a pair: the constant M_{eff} it holds, and the number of binades over which it holds it. For the TNF ladder those are $(M, 2 \cdot 3^{E_t-1})$ by construction; for a tapered format neither is constant, which is why a ladder makes a better ruler than a taper.*

4.1 The diagnostic across a catalogue

Applying Theorem 2 to every format for which a published oracle was available — 51 in all, spanning IEEE, bfloat, posit, takum, LNS, MXFP, VAX, IBM hexadecimal, Cray, PDP-11, x87 and Microsoft binary, plus both GoldenFloat ladders — produced three observations we did not anticipate.

Declared widths are recovered across three orders of magnitude of width. The GF ladder returns 8.97 at GF16 against a declared 9, 78.03 at GF128 against 78, and 631.99 at GF1024 against 632. The same holds for double-double (106.68), quad-double (215.90) and x87’s 80-bit format (63.04). A relation that reconstructs a declared parameter from round-trip error alone, at widths from 8 to 1006 bits, is doing something other than describing the format it was derived on.

The taper rate is a constant of the family — but only posit has a rate. posit16 sheds 0.261 bits of significand per binade and posit32 sheds 0.247, against the $2^{-es} = 0.25$ that Theorem 3 below requires; posit64, whose oracle carries the pre-standard $es = 3$, sheds 0.123 against a required 0.125. Three formats, three exact confirmations. We previously reported a comparable per-binade rate for takum, and that was a category error on our part: takum’s staircase is logarithmic, so fitting a line to it produces a number that depends on the interval fitted and describes nothing. Its correct constant is one bit per *doubling* of $|e|$, and against that the four takum and three tekum rungs measure 0.83 to 1.09.

Range exhaustion is distinguishable from tapering, and both from neither. The distinction is not optional: our own first pass placed the measurement bands at fixed positions rather than inside each format’s measured range, and read the saturation of every narrow format as a taper — classifying TNF16, a fixed-field format, as geometrically tapered. Measuring each format’s usable range first and then placing the bands inside it removes the artefact entirely. GF14 reads flat at 8.05 across its range and collapses only below its subnormal boundary, which is not a defect; the instrument says which of the two it is, where an accuracy table alone does not.

4.2 Tapering has a functional form, and it follows from the encoding

Section 4.1 reported a taper rate per family. That was half right, and the half that was wrong is instructive: a rate presupposes a linear taper, and only one of the two families has one.

Theorem 3 (Posit taper, exact). *For a posit of n bits with exponent field es , the regime is a run of k like bits carrying a scale of $used^k = 2^{2^{es}k}$ and occupying $k + 1$ bits. Reading k out of the encoder for 2^e confirms $k = \lfloor |e|/2^{es} \rfloor + 1$ for every e tested. The significand therefore loses*

$$\Delta M_{posit}(e) = \left\lfloor \frac{|e|}{2^{es}} \right\rfloor \quad \text{bits: a staircase, arithmetic in } |e|.$$

The Posit Standard (2022) fixes $es = 2$ at every precision [6], so the step is one bit per four binades at all widths — which is why the loss appeared as a family constant of 0.25 bits per binade. Measured slopes of -0.2497 and -0.2505 at posit16 and posit32 are the staircase seen through a linear fit.

Theorem 4 (Takum taper, exact). *Takum spends a fixed 3-bit regime field holding r , after which the characteristic occupies r bits. Reading r out of the encoder for 2^e gives $r = \lfloor \log_2(|e| + 1) \rfloor$ exactly, at 16, 32 and 64 bits alike. The significand therefore loses*

$$\Delta M_{takum}(e) = \lfloor \log_2(|e| + 1) \rfloor \quad \text{bits: a staircase, geometric in } |e|.$$

Corollary 2 (The two families differ in kind, not degree). *Within its range a fixed-field format loses nothing; a posit loses one bit of significand per 2^{es} binades without limit; a takum loses one bit per doubling of $|e|$. Tabulating ΔM :*

$ e $	fixed-field	posit ($es=2$)	takum
30	0	7	4
100	0	25	6
1,000	0	250	9
10,000	0	2,500	13

At $|e| = 1,000$ a posit would need 250 bits of significand merely to pay for its regime, which is why 16- and 32-bit posits cannot reach there at all. This is the mechanism behind takum’s stated improvement on posits [1], stated as a quantity.

Table 3: The catalogue by staircase form. The parameter is the taper constant in the units natural to each form. Measured against each format’s own usable range.

form	M_{eff} against $ e $	count	members and constants
constant	flat	38	IEEE, bfloat, VAX, Cray, x87, GF, TNF
wobble	periodic, no trend	3	IBM hexadecimal, $\log_2 16 = 4$ bits
arithmetic	-2^{-es} per binade	5	posit16/32 0.25, posit64 0.125
geometric	-1 per doubling	8	takum $\times 4$, tekum $\times 3$

Corollary 3 (A rate is not always the right summary). *Both laws are staircases, so a fitted slope is a summary and not the thing itself, and it is a faithful summary only when the staircase is arithmetic. An earlier draft of this paper reported -0.117 bits per binade for takum16 over $|e| \in [0, 30)$; the same format over $[0, 60)$ gives a smaller figure for no reason but the window, and the smooth-log fit of $c = 0.75$ understates the true step of one bit per doubling for the same reason. Identify the form before choosing a statistic — and prefer reading the encoder to fitting its output.*

Corollary 4 (A three-way taxonomy, recovered without knowing the encoding). *The diagnostic separates formats by the form of $M(e)$: constant (fixed-field, TNF and GF by construction), linear in $|e|$ (posit), or logarithmic in $|e|$ (takum). Which form fits is decided by the data, not by reading the specification — and the fitted coefficients then recover the encoding parameters, 2^{-es} for posit and the characteristic width for takum.*

What the literature describes qualitatively — more precision near unity, less at the extremes [6] — turns out to have an exact form recoverable two ways that agree: from the encoder, by reading the regime field, and from round-trip error alone, by the diagnostic of Theorem 2. The second route needs no access to the specification, which is what makes it useful on a catalogue.

4.3 Four staircase forms, and only four

Classifying by the *form* of the staircase rather than by a fitted slope resolves the catalogue into four classes. Of the 83 entries, 11 are integer or fixed-point formats with no exponent field; of the remaining 72, 18 span fewer than six binades and cannot be classified from round-trip error at all, and one requires probes wider than we generated.

The fourth class was not in our scheme until the catalogue produced it, and it has a closed form.

Theorem 5 (Precision wobble). *Let a format scale by r^e with a significand quantised uniformly over a scale interval spanning a factor r . Its relative error varies by a factor of r across that interval, so M_{eff} oscillates with amplitude exactly $\log_2 r$ bits and has no trend in $|e|$.*

Proof. The quantisation step is constant across the interval and the relative error is that step divided by the significand, which ranges over a factor r . Taking $-\log_2$ of the ratio of the extremes gives $\log_2 r$. Since the interval repeats identically at every e , the oscillation is periodic and carries no trend. $\square$ $\square$

Theorem 5 reproduces a result sixty years older than this work. IBM System/360 hexadecimal is documented as delivering between $4p-3$ and $4p$ significant bits, a spread of three [12, 13]; the integer count is the floor of the continuous quantity, so the theorem’s $\log_2 16 = 4$ and the textbook’s 3 are the same statement. Measured over eight sub-intervals — which by averaging within each bin predicts a reduced spread of 2.96 rather than the full 4 — `ibm_hfp32` and `ibm_hfp64` give 3.13 and 3.04, while `binary32` and `gf32` give 0.85 and 0.93 against a predicted

Table 4: Round-trip mean relative error. Bins are powers of two, not decades.

magnitude	GF16 (φ)	TNF16	tekum16	TNF16 vs tekum16
$ e < 8$	3.43e−4	3.56e−4	3.27e−4	0.92× (tie)
$ e 8\text{--}20$	3.57e−4	3.52e−4	1.00e−3	2.84 ×
$ e 20\text{--}38$	6.98e−3 (478 clip)	3.53e−4	1.95e−3	5.53 ×

0.87. Binary formats wobble by one bit; this is the familiar 1-bit wobble of every radix-2 float, and it is the smallest wobble any integer radix admits. Theorem 5 therefore supplies a second and independent argument for TNF’s binary scale, alongside the 0.331 positions of Theorem 10 below: a radix-3 scale would wobble by 1.585 bits instead of 1.

The taxonomy also settles why a format must choose between the first class and the last two, and hence why this work is a ladder rather than a format.

Theorem 6 (Range–precision dichotomy). *Fix a width N . If a format’s exponent takes unboundedly many values, then its exponent codewords have unbounded length, and $M_{\text{eff}}(e) \rightarrow 0$ as $|e| \rightarrow \infty$. Consequently no fixed-width format has both unbounded range and precision bounded below by a positive constant.*

Proof. The exponent codewords form a prefix code over a binary alphabet, so by Kraft’s inequality $\sum_e 2^{-\ell(e)} \leq 1$. If the sum has infinitely many terms then $\ell(e) \rightarrow \infty$. The payload available to the significand is $N - 1 - \ell(e)$, so it falls to zero, and by Theorem 1 so does M_{eff} . $\square$ $\square$

Corollary 5 (Why a ladder). *Theorem 6 makes constant precision and unbounded range incompatible at any fixed width, so a format that wants both must obtain the second by varying N . A ladder of fixed-field rungs does exactly this: each rung holds its precision flat across its whole range, and the ladder as a whole covers any range required. The tapered families make the opposite choice — one rung, unbounded range, precision decaying at the rate of their regime code, arithmetic for posit and geometric for takum. Neither choice is better in the abstract; what the dichotomy establishes is that they are the only two available.*

Read against Corollary 5, the three tapered constants in Table 3 become a statement about codes rather than about formats. A unary regime costs $\ell(e) \propto |e|$ and yields the arithmetic staircase; a length-prefixed binary regime costs $\ell(e) \propto \log_2 |e|$ and yields the geometric one; a fixed field costs $\ell(e) = \text{const}$ and yields no staircase at the price of bounded support. The staircase form is the growth rate of the exponent’s codeword length, and nothing else. This also names the gap in the table: between $\ell \propto |e|$ and $\ell \propto \log_2 |e|$ no catalogued format sits. Section 12 builds the codes and asks what is in the gap; the answer corrects a suggestion we made before building them.

5 Accuracy

Section 3 predicts a rung’s error from its mantissa width without measuring it. This section measures it, against the format the ladder was built to be compared with. The metric is mean relative error over an encode–decode round trip, 6,000 values spanning $2^{-38} \dots 2^{38}$ with random sign, binned by binary-exponent magnitude $|e|$. Table 4 and Figure 1.

On the axis. The bins are powers of two. Read as *decades* the far column is not a win at all: TNF16’s exponent reaches ± 40 in powers of two, roughly ± 12 decades, so beyond that it overflows while an unbounded regime keeps working — measured, 1,458 of the far-bin values clip under that reading. We state the axis explicitly because an earlier internal note did not, and the omission points a reader at the one interpretation under which the result appears fabricated.




Figure 1: Accuracy by binary-exponent magnitude. GF16’s 6-bit exponent clips 478 of 2,857 values in the far bin; TNF16’s balanced-ternary exponent does not.

Table 5: Mean relative round-trip error by binary-exponent magnitude, with overflow counts in parentheses. 6,000 values, $2^{-38} \dots 2^{38}$, random sign, one seed.

format	$ e < 8$	$ e 8-20$	$ e 20-38$
TNF16	3.56e−4	3.52e−4	3.53e−4
GF16 (ϕ)	3.56e−4	3.52e−4	6.76e−3 (478)
takum16 / tekum16	3.27e−4	1.00e−3	1.95e−3
posit16	1.36e−4	8.31e−4	1.43e−2
IEEE binary16	1.76e−4	1.28e−3 (299)	1.57e−1 (2479)
bfloat16	1.45e−3	1.38e−3	1.41e−3
LNS16	8.39e−4 (1324)	7.63e−4 (1808)	6.66e−4 (2849)

6 The 16-bit field

Comparing against one incumbent invites the objection that the incumbent was chosen. Table 5 runs every 16-bit-class format for which this repository ships a published oracle [5] through the same workload and the same metric. Values that overflowed are counted separately rather than folded into the mean, since a format should not look accurate for having discarded the hard cases.

Three readings, including the one that goes against us.

TNF16 is flat, and alone in being flat without clipping. Its error is 3.5e^{-4} at every magnitude in the workload and it discards nothing. bfloat16 is also flat and clip-free, at roughly four times the error. Every other entrant either degrades with magnitude or overflows: binary16 does both, losing three orders of magnitude and discarding 2,479 values; GF16’s 6-bit exponent clips 478; LNS16 clips throughout.

At 32 bits the binary formats win outright. Running the same workload on the 32-bit class: GF32 is flat at 3.4e^{-7} with no clipping, but IEEE binary32 is flat at 2.1e^{-8} , and posit32 and takum32 are better still near unity (2.1e^{-9} and 4.9e^{-9}). Uniformity stops being

Table 6: The whole ladder, measured in exact rationals. From TNF16 upwards the error is flat in magnitude and its exponent roughly doubles per rung. TNF4 and TNF8 are shown at their limits: the workload spans 76 binary decades and those rungs hold 2 and 8, so most of it falls outside them — reported rather than omitted, since a ladder’s lower rungs are where its range trade is visible.

rung	decades	$ e < 8$	$ e 8\text{--}20$	$ e 20\text{--}38$
TNF4	2	beyond range for this workload		
TNF8	8	1.22e−2	beyond range	
TNF16	24	3.45e−4	3.69e−4	3.66e−4
TNF32	219	5.27e−9	5.63e−9	5.58e−9
TNF64	658	4.33e−17	4.22e−17	4.12e−17
TNF128	1,975	4.25e−36	4.55e−36	4.51e−36
TNF256	5,925	2.76e−74	2.69e−74	2.63e−74
TNF512	17,775	4.32e−151	4.62e−151	4.58e−151
TNF1024	53,326	2.84e−304	2.77e−304	2.71e−304

the deciding property once there are enough bits that nothing clips anyway. The TNF argument is a 16-bit-class argument, and we do not extend it.

Near unity we lose. posit16 at 1.36e−4 and binary16 at 1.76e−4 are more accurate than TNF16 there, by a factor of two and a half and of two. A tapered format spends its bits where the values usually are, and near unity that is the right bet. TNF16 buys uniformity instead, and the trade only pays for a workload that leaves the neighbourhood of one.

takum16 and tekum16 are the same format here, exhaustively. The two oracles produce identical figures on every bin, so we checked the whole 16-bit code space: all 65,536 codes decode identically, with zero differences, although the two implementations differ by 242 lines. The reconstruction our tekum oracle implements *is* takum16. Every comparison in this paper should therefore be read as against takum16, and we relabel it as such. That is not a weaker anchor — takum is published, has a public VHDL codec, and is the parent of the format we set out to compare against.

7 Hardware

Synthesis with Yosys 0.65 [7], place and route with nextpnr-xilinx (1743d0f) [8] on a Xilinx XC7A200T, hard multipliers disabled so the figures describe fabric alone. Bitstream generation uses Project X-Ray. No vendor tool is involved at any step, which is what makes the numbers checkable by a reader without a licence.

For context rather than as a controlled comparison: a published Altera ALTFP_MUL on a Cyclone IV reports 119–132 MHz at 6–10 cycles of latency, using 832–1041 logic elements *and* 18 embedded multipliers, and posit multiplier studies report 95–572 LUTs with 4.28–15.55 ns of logic delay. Different devices and different toolchains; we draw no equivalence from them beyond the observation that TNF16’s figures sit inside that envelope on all four axes at once, on a fabric that gives the format no ternary advantage.

Isolating the pipeline register: the combinational multiplier closes at 81.35 MHz, the two-stage version at 147.32 MHz. The cut is between the significand product and the renormalisation — a 10×10 multiply, then a carry test, a saturating exponent add and a bit select.




Figure 2: The same data. Flatness across the three magnitude bands is the property a fixed-field format is built for; the shaded rungs are those the workload overruns.

Table 7: TNF multiplier, post-route, one harness across the whole lineup so the rows are comparable. Latency one cycle, throughput one result per cycle.

rung	LUTs	DSP48	$F_{\max}$
TNF4	12	0	161.11 MHz
TNF8	50	0	153.23 MHz
TNF16	212	0	131.73 MHz
TNF32	1,477	0	83.27 MHz
TNF64	6,140	0	53.26 MHz
TNF128	10,029	0	33.23 MHz

8 Two width defects, and why they generalise

Section 7 reports what the datapath costs once every bus is derived from the format’s own parameters. It was not measured that way first. The two defects below — one of area, one of correctness — are the same mistake, a width chosen by hand, and they are reported because that mistake is not specific to TNF.

8.1 Interface width dominated the arithmetic

Our first synthesisable realisation declared every port and the significand product 32 bits wide. Nothing in TNF16 is 32 bits: the mantissa field is 9, so $1+M$ is 10 and their product is exactly 20; the exponent offset never exceeds $\text{OFFSET_MAX} = 80$, which is 7. Synthesis built a 32×32 multiplier and a 32-bit comparison tree and charged for it (Figure 5): 1,179 LUTs, or three DSP48 blocks, against 219 LUTs and none once the buses match the values they carry. The arithmetic is unchanged.




Figure 3: The same data. Scissors mark values that overflowed.

Table 8: The specified ladder. E_t is the trit budget, M the mantissa width, range in decades. Rungs above TNF128 exceed a single Artix-7 fabric and are oracle- or ASIC-scale; they are specified and covered by the reference oracle, not measured here.

rung	E_t	M	decades	rung	E_t	M	decades
TNF4	2	1	2.4	TNF64	7	52	658
TNF8	3	4	8	TNF128	8	115	1,975
TNF16	4	9	24	TNF256	9	242	5,925
TNF32	6	25	73	TNF512	10	497	17,775
				TNF1024	11	1006	53,326

8.2 The same width decided correctness

At TNF32 the significand product needs 52 bits. The 32-bit wire truncates, and the module returns wrong results while appearing to support the rung, since its own documentation names TNF32 as a supported configuration. Measured against a 64-bit reference: 1,995,730 mismatches in 2,128,964 combinations.

Deriving every width from the format parameters removes both problems at once. We report this because the pattern is not specific to TNF: in a parametric arithmetic core, a width that is merely *large enough for the configuration that was tested* is a silent-truncation trap at the next one, and it will not be caught by a testbench written against the tested configuration.

9 Equivalence

Every change above is a claim that timing or width changed and arithmetic did not. Each was checked rather than argued.

The TNF16 sweep is exhaustive in the sense that matters here: the mantissa space is covered in full at offset pairs exercising underflow, mid-range and saturation, then the offsets are covered in full at mantissas that do and do not carry — every path through the carry, the saturation and the underflow clamp. The last row is the point of the section: the width-derived module agrees with the reference that is correct and disagrees with the one that truncates.




Figure 4: Area and achieved frequency across the ladder.

Table 9: Equivalence evidence. The reference for each rung is the module already trusted for that rung.

check	combinations / cycles	mismatches
width-derived vs original, TNF16 (sweep)	321,156	0
width-derived vs original, TNF4	374	0
width-derived vs original, TNF8	3,716	0
width-derived vs original, TNF16	77,444	0
width-derived vs 64-bit reference, TNF32	300,000	0
pipelined vs combinational	199,994	0
original 32-bit module vs 64-bit ref, TNF32	2,128,964	1,995,730

10 What the law says about improving the ladder

Theorem 1 makes the design space arithmetic rather than a matter of taste. Error scales as $2^{-(M+1)}$ and range as 3^{E_t} , so at a fixed width the only decision is how many bits the exponent field is allowed to consume — and that is decided by how efficiently 3^{E_t} values pack into $\lceil E_t \log_2 3 \rceil$ bits.

The width convention counts positions, and one rung breaks it. A rung’s width is its position count with a trit occupying one position, so the layout is $1 + E_t + M = N$. Three of the four specified rungs satisfy it exactly: TNF4 is $1 + 2 + 1$, TNF8 is $1 + 3 + 4$, TNF32 is $1 + 6 + 25$. TNF16 is $1 + 4 + 9 = 14$, two positions short of its name, and neither specification file gives a reason. The likely history is visible in the parameters: TNF16 inherits GF16’s φ -optimal $M = 9$ and replaces GF16’s six-bit exponent with four trits, which buys range — $3^4 = 81$ steps against $2^6 = 64$ — and frees two positions that were never re-spent. It is the one rung derived from its parent rather than from the width rule.

This also settles what a rung costs on a binary fabric, which is a separate question from its width: a trit stored as a two-bit code makes TNF16 eighteen bits, and the offset stored as a plain binary integer makes it seventeen. Those are fabric costs, not format widths, and confusing the two is what made the ladder appear not to fit.




Figure 5: The same arithmetic at three interface widths.

Table 10: Packing efficiency of the exponent field. The theoretical floor is $\log_2 3 = 1.585$ bits per trit.

E_t	values 3^{E_t}	bits	capacity	efficiency
3	27	5	32	84.4%
4	81	7	128	63.3%
5	243	8	256	94.9%
6	729	10	1024	71.2%
7	2187	12	4096	53.4%
8	6561	13	8192	80.1%
10	59049	16	65536	90.1%

The slack is not confined to one rung. Applying $1 + E_t + M = N$ to every specified rung separates the two documents that define the ladder. The four rungs of the machine-readable specification obey it at TNF4, TNF8 and TNF32; the five upper rungs, which exist only in a research note, miss it by four to six positions each, and the note and the specification disagree with each other about TNF32.

No single alternative convention explains the upper rungs either: reading the exponent as $\lfloor E_t \log_2 3 \rfloor$ bits fits TNF16, TNF64, TNF128 and TNF1024 exactly but overshoots TNF256 and TNF512 by one and misses TNF32 entirely. We conclude that the upper ladder was tabulated rung by rung rather than derived, and that reconciling it costs nothing at the rungs that are already right.

What the two free positions were worth, and where they went. Because the exponent cancels in Theorem 1, the menu at sixteen positions could be read off before measuring, and the measurement confirmed it to four significant figures:

The last row is the law’s own consistency check: adding two trits of exponent while holding M leaves the error bit-identical and multiplies the range ninefold.

Table 11: The width rule across the ladder. k is the shortfall $N - (1 + E_t + M)$. The reconciled column is $M = N - 1 - E_t$, which changes nothing at the three rungs that already obey the rule.

rung	source	E_t	M	$1+E_t+M$	k	reconciled M
TNF4	spec	2	1	4	0	1
TNF8	spec	3	4	8	0	4
TNF16	spec	4	9	14	2	11
TNF32	spec	6	25	32	0	25
TNF32	note	5	21	27	5	—
TNF64	note	7	52	60	4	56
TNF128	note	8	115	124	4	119
TNF256	note	9	242	252	4	246
TNF512	note	10	497	508	4	501
TNF1024	note	11	1006	1018	6	1012

Table 12: The three ways to spend TNF16’s two unallocated positions. Error and range trade against each other exactly as Theorem 1 predicts; nothing in any row is worse than the unfilled rung.

variant	positions	mean rel. error	decades	vs. takum16 (far)
unfilled, $E_t=4$, $M=9$	14	3.37e−4	24	5.46×
$E_t=4$, $M=11$ (adopted)	16	8.63e−5	24	22.68×
$E_t=5$, $M=10$	16	1.69e−4	73	—
$E_t=6$, $M=9$	16	3.37e−4	219	—

The positions were spent on mantissa, and the reason is a measurement. Corollary 10 forbids recommending a budget without naming a workload, so we named one. At the 16-bit class the extra range is not consumed: quantised-inference tensors span roughly 2^{-14} to 2^{14} , and $E_t = 4$ already clips nothing at ± 40 binades — zero clips in 2000 values. $E_t = 6$ would have bought ± 364 binades at bit-identical error, which is range nobody spends at sixteen bits. So $M = 11$: the error falls by exactly 4.00, and against takum16 the rung moves from $0.92 \times / 2.88 \times / 5.49 \times$ to $3.70 \times / 11.21 \times / 22.68 \times$, no longer losing the near bin. In silicon it costs 443 LUTs against 372, a 19% increase, at 136.44 MHz against 131.73 — no frequency penalty.

All three ratios in this section and in Figure 1 come from one script, `gen_figures.py`, which computes them from the oracles at seed 20260809 rather than transcribing them. They were hardcoded until 2026-08-09 and had drifted: the figure still carried $M = 9$ and the label `tekum16` after the rung moved and the oracle was shown to decode identically to takum. A figure that disagrees with the table it illustrates is the same defect class as a gate that cannot go red.

Conformance was re-run in full rather than assumed: round-trip with sign 2000/2000, encoding monotone with zero inversions, $\pm$ symmetry with zero mismatches in 500, both range boundaries decoding. A `width_rule` test now asserts $1 + E_t + M = N$ in the specification itself, so the position count is checked rather than remembered. All nine rungs satisfy it.

11 Why ternary, and exactly how much

The case for a ternary exponent is usually made by citing radix economy and left there. We state it, then bound it in two directions — what a binary fabric takes back, and what the radix

Table 13: Utilisation $\rho(E_t)$ of the binary field holding an E_t -trit exponent. The ladder’s present budgets, $E_t \in \{4, 7\}$, are the two worst entries in the range.

E_t	2	3	4	5	6	7	8	9	10	11
ρ	56.2	84.4	63.3	94.9	71.2	53.4	80.1	60.1	90.1	67.6

of the *scale* is worth as distinct from the radix of the *exponent* — because both bounds turn out to be decisive for the design and neither is folklore.

Theorem 7 (Radix economy). *Representing R distinct values in radix b costs $\log_b R = b \ln R / \ln b$ state-slots. The factor $b / \ln b$ is minimised at $b = e$; among integers it is $3 / \ln 3 = 2.731$ against $2 / \ln 2 = 2.885$, so a ternary field is 5.3% more economical than a binary one.*

This is the classical argument [9, 10] and it is what motivates a ternary exponent. It is also stated in a currency — state-slots — that no binary fabric pays in. The next theorem says what happens when it does.

Theorem 8 (No free range on a binary fabric). *An exponent field of E_t trits spans 3^{E_t} values. Stored as a binary integer it occupies $\lceil E_t \log_2 3 \rceil$ bits, a field spanning $2^{\lceil E_t \log_2 3 \rceil} \geq 3^{E_t}$ values. Hence on a binary fabric a ternary exponent never spans more values per bit than a binary one, with equality only at $E_t = 0$. Writing $\rho(E_t) = 3^{E_t} / 2^{\lceil E_t \log_2 3 \rceil}$ for the utilisation, $\rho \in (1/2, 1]$ and ρ does not converge.*

Run as an optimisation rather than an inequality, the same statement is sharper than it looks. Asking which exponent encoding a binary fabric prefers at $(N, R) = (16, 24)$, $(32, 219)$ and $(64, 658)$ returns the same mantissa width for both — $M = 10, 23$ and 53 respectively. The ternary encoding does not lose on a binary fabric; at these widths the ceilings absorb the difference and it exactly ties, which is the strongest form of “never more values per bit”.

Theorem 8 is the reason the ladder’s rung widths must be read as position counts and not bit counts, and it is where our own first reading of the ladder went wrong: measuring what a trit costs when packed into bits and calling the result the format’s width confuses a fabric cost with a format property. The independently published balanced-ternary float Ternary27 [11] counts the same way — 2 type trits, 1 sign trit, 5 exponent trits and 19 significand trits, exactly 27 — which is external corroboration that position counting is the convention in this family rather than a reading of our own.

Theorem 9 (Unallocated positions are a pure loss). *Let a rung of N positions carry a sign, E_t exponent trits and M mantissa bits with $1 + E_t + M = N - k$, $k \geq 0$. Then relative to the rung with the same E_t and $M + k$, its expected relative error is larger by exactly 2^k and its dynamic range is identical.*

Proof. By Theorem 1 the expected relative error is $\frac{1}{2} \mathbb{E}[1/s] 2^{-(M+1)}$, in which the exponent does not appear; replacing M by $M + k$ multiplies it by 2^{-k} . The dynamic range is 3^{E_t} binades, a function of E_t alone, so it is unchanged. Nothing else in the layout depends on M . $\square$ $\square$

Theorem 9 converts the slack of Section 10 into a number without any measurement, and the measurement then confirms it. Against a workload built in exact arithmetic — necessary above TNF32, where a double-precision probe measures its own resolution rather than the format’s — the predicted and observed factors are:

The final question is one the ladder answers implicitly and, as far as we can find, no one has stated: TNF encodes its exponent in radix 3 but scales by 2^e , whereas Ternary27 scales by 3^e . These are different choices and they are separable.

Table 14: Unallocated positions per rung, the loss Theorem 9 predicts, and the loss measured against the reconciled rung $M = N - 1 - E_t$. Exact-arithmetic workload, 800 probes per rung.

rung	k	predicted	measured
TNF16	2	$4\times$	$4.03\times$
TNF64	4	$16\times$	$15.43\times$
TNF128	4	$16\times$	$15.99\times$
TNF256	4	$16\times$	$15.24\times$

Theorem 10 (The scale radix, and why 2 beats 3). *Let a format scale by r^e with a normalised significand uniformly quantised over $[1, r)$ by M bits. Then its expected relative error on a log-uniform workload is $\kappa(r) 2^{-M}$ with*

$$\kappa(r) = \frac{(r-1)^2}{r \ln r}, \quad \kappa(2) = 0.72135, \quad \kappa(3) = 1.21365.$$

Holding the field layout fixed, moving the scale from $r = 2$ to $r = 3$ therefore multiplies the error by $\kappa(3)/\kappa(2) = 1.6825$ while multiplying the dynamic range by $\log_2 3 = 1.5850$. Measured in positions, it costs $\log_2 1.6825 = 0.7506$ of precision and buys $\log_3 1.5850 = 0.4192$ of range: a net loss of 0.3314 positions per number. A power-of-two scale strictly dominates a power-of-three scale at equal width.

Proof. The quantisation step over $[1, r)$ is $(r-1)2^{-M}$ and the mean absolute rounding error is half of it. On a log-uniform workload the significand has density $1/(s \ln r)$ on $[1, r)$, so $\mathbb{E}[1/s] = \frac{1}{\ln r} \int_1^r s^{-2} ds = \frac{1-1/r}{\ln r}$. Multiplying gives $\frac{1}{2}(r-1)2^{-M} \cdot \frac{1-1/r}{\ln r}$, and $\frac{1}{2}(r-1)(1-1/r) = (r-1)^2/2r$. Absorbing the factor $\frac{1}{2}$ into the $2^{-(M+1)}$ convention of Theorem 1 yields κ as stated. For the range, E_t trits span 3^{E_t} steps of $\log_2 r$ binades each. One mantissa bit halves the error and one exponent trit triples the range, which fixes the exchange rate between the two currencies and gives the position figures. $\square$ $\square$

Direct measurement at two rungs, encoding and decoding through both scales in exact arithmetic, gives error ratios of 1.675 at $(E_t, M) = (4, 11)$ and 1.690 at $(6, 25)$ against the predicted 1.6825, and range ratios of 1.583 and 1.585 against the predicted 1.5850.

Corollary 6 (What TNF is, precisely). *TNF is not a ternary float in the sense of Ternary27; it is a binary-radix float whose exponent is encoded in balanced ternary. Theorem 10 says that this is the better of the two available choices by 0.331 positions per number, and Theorem 8 says that the ternary encoding pays only on a ternary fabric. The two together delimit the format exactly: the radix choice is right on any fabric, and the encoding choice is right on one.*

12 The regime field is a universal code

Theorem 6 identifies the exponent’s encoding as a prefix code. That identification is worth making explicit, because it places the taxonomy of Section 4.3 inside a hierarchy that predates every format in the catalogue by decades [14].

Theorem 11 (The staircase is the codeword-length function). *Let a format of N positions encode its exponent with a prefix code assigning e a codeword of length $\ell(e)$. Then $M_{\text{eff}}(e) = N - 1 - \ell(e)$, and the staircase form is the growth rate of ℓ and nothing else.*

The catalogue’s three tapering classes are then three points in the classical hierarchy: a unary regime, $\ell \propto |e|$, gives posit’s arithmetic staircase; a length-prefixed binary regime, $\ell \propto \log_2 |e|$, gives takum’s geometric one; a fixed field, $\ell = \text{const}$, gives no staircase at the price of bounded support. Elias’ γ , δ and ω codes occupy the logarithmic class at successively finer asymptotics.

Table 15: M_{eff} by $|e|$ and usable range for three regime codes at $N = 16$, each normalised to satisfy Kraft. The square-root code is a genuine Pareto point: more precision than the logarithmic code throughout the core, more range than the unary one.

regime	$ e =1$	4	16	64	256	binades
unary, $\ell \propto e $ (posit class)	11	10	7	-5	-53	43
$\ell = \lceil \sqrt{ e } \rceil$	11	10	8	4	-4	121
$\ell \propto \log_2 e $ (takum class)	9	7	5	3	1	511

Theorem 12 (Logarithmic is the floor). *For any prefix code on the integers, $\ell(e) \geq \log_2 |e|$ for infinitely many e . Hence no format of unbounded range tapers more gently than one bit per doubling of $|e|$, asymptotically.*

Proof. Kraft’s inequality gives $\sum_e 2^{-\ell(e)} \leq 1$. If $\ell(e) < \log_2 |e|$ for all but finitely many e , the tail of the sum exceeds $\sum_e 1/|e|$, which diverges. $\square$ $\square$

Theorem 12 is worth stating plainly because it bounds this work as much as its competitors: takum’s regime is asymptotically optimal, and no format — TNF included — can beat it over unbounded range. What TNF does instead is decline the unbounded range, which Theorem 6 shows is the only other option. The measured advantages of Section 5 are therefore advantages *inside TNF’s range*, which is where the taper has already cost takum the bits, and they should never be read as a claim beyond it.

Theorem 13 (The regime’s radix does not matter). *A regime spending one digit per multiplication of $|e|$ by r , over an alphabet of r symbols, costs $\log_2 r \cdot \log_r |e| = \log_2 |e|$ bits, independent of r . The staircase form is invariant to the regime’s radix.*

Theorem 13 retracts a construction we proposed in an earlier draft of this work. Having observed the gap between the arithmetic and geometric classes, we suggested a regime spending one trit per tripling of $|e|$ as the natural way for a ternary format to occupy it. It does not occupy it: building both codes and normalising each to satisfy Kraft, the two codeword-length functions differ by an additive constant of exactly 1 at every $|e|$ from 1 to 4096. One trit per tripling is takum’s staircase with a different constant, not a third form. A ternary fabric changes what the regime costs to decode; it does not change what it costs to encode.

The gap is real, occupiable, and empty for a reason we can only conjecture. Growth rates strictly between $\log_2 |e|$ and $|e|$ do exist and are Kraft-feasible — $\ell(e) = \lceil \sqrt{|e|} \rceil$ is one — and such a code is not dominated. At $N = 16$, normalised so that all three are legal prefix codes:

For reference, a fixed field reaching the same 121 binades as the square-root code spends 8 bits on the exponent and holds 7 bits flat everywhere — better than the square-root code beyond $|e| \approx 30$ and worse below it, which is the taper trade in miniature.

So the gap is not empty because its occupants are dominated; they are not. Our first explanation was decode cost, and building the hardware refutes it.

12.1 What a regime costs in silicon

We wrote three regime codecs to one interface — a 16-bit word in, an exponent and a mantissa shift out, and the mirror encoder — and synthesised each under a single harness on XC7A200T with Yosys and nextpnr-xilinx, hard multipliers disabled. These are structural models of the three classes, not reproductions of the published posit or takum codecs; what they compare

Table 16: Regime codec cost by class, post-route on XC7A200T under one harness, no DSP. Figures are net of the harness, which costs 24 LUTs and was synthesised separately and subtracted; the round-trip column is decode plus encode.

class	decode		encode		round trip
	LUTs	$F_{\max}$	LUTs	$F_{\max}$	LUTs
unary, $\ell \propto e $ (posit)	379	36.8 MHz	59	211.2 MHz	438
$\ell \propto \log_2 e $ (takum)	6	805.8 MHz	34	285.1 MHz	40
$\ell = \lceil \sqrt{ e } \rceil$	393	33.1 MHz	123	96.4 MHz	516
fixed field (GF, TNF)	0	—	0	—	0

is a run-scanned regime against a length-prefixed one, which is the structural difference the taxonomy names.

The conjecture said the square-root code would be too expensive to be worth building. It is not: it costs 18% more than the unary regime round trip, and the unary regime ships in production. What the measurement does show is that the expense lies somewhere else than we assumed, and this has a closed form.

Theorem 14 (Decode cost is set by the scan, not by the staircase). *A regime whose codeword length is delimited by a run must examine up to $N - 1$ positions to decode, giving cost $\Theta(N)$. A regime whose length is carried in a fixed field of b bits requires no scan and a shift of at most $2^b - 1$, giving cost $\Theta(1)$ in N . The two costs are independent of the growth rate of ℓ , and therefore of the staircase form.*

Theorem 14 predicts the table exactly. The unary and square-root codes share the same 15-position scan and cost 403 and 417 LUTs, a difference of 3.5% that is the squarer and nothing else; the length-prefixed code has no scan and costs 30. The square-root code is geometrically intermediate and structurally identical to the unary one, which is the point: *the staircase form and the decode cost are independent axes*, and a taxonomy of one does not predict the other.

Corollary 7 (The logarithmic regime is not a trade). *Against the unary regime the length-prefixed one is better on every axis measured: asymptotically optimal by Theorem 12, 11× smaller round trip, and 22× faster to decode. The published bounded-regime posit variants [15], which cap the regime field precisely to avoid the scan, are the same conclusion reached from the other direction.*

Corollary 8 (The measured price of unbounded range). *A fixed field has no regime codec at all. The dichotomy of Theorem 6 therefore has a hardware statement: unbounded range costs 40 LUTs at best and 438 at worst on this device, against 0 for a format that declines it. Section 13 converts those LUTs into the same currency as the accuracy results, which is the only way to say whether they matter.*

We are left without an explanation for the empty gap. It is not domination, since the square-root code is Pareto-undominated; and it is not cost, since it is within 16% of a regime that ships. Both of our explanations are now measured and both are wrong, and we record that rather than replace it with a third.

13 Converting silicon into precision

Corollary 8 gives a regime’s cost in LUTs and Section 5 gives an accuracy ratio, and the two cannot be compared until they are in one currency. They can be: at a fixed class, a mantissa bit has a marginal area, and by Theorem 1 it has a known value in accuracy.

Table 17: Regime area converted into mantissa bits through Theorem 15. The unary regime costs, in silicon alone, more than TNF16’s entire mantissa.

regime	LUTs	bits at 16-bit class	bits at 32-bit class
unary (posit class)	438	8.98	3.94
length-prefixed (takum class)	40	0.82	0.36

Theorem 15 (Area and precision are commensurable). *Let $\lambda = dA/dM$ be the marginal area of a mantissa bit at a given class. A regime codec costing C LUTs is then equivalent to C/λ mantissa bits of silicon, and by Theorem 1 to a factor $2^{C/\lambda}$ in accuracy at equal area.*

Measuring λ directly — the same width-derived multiplier synthesised at $M = 7, 9, \dots, 25$ under one harness — gives 285, 372, 443, 567, 696, 870, 1238 and 1683 LUTs. The marginal cost is not constant: it rises from 35.5 LUTs per bit between $M=9$ and $M=11$ to 111.2 between $M=21$ and $M=25$. Taking the local value at each class, $\lambda = 48.8$ at the 16-bit class and 111.2 at the 32-bit one:

13.1 How area grows along the ladder

Theorem 15 needs λ , and λ is a derivative of an area law we can state. Fitting $A(M) = cM^\alpha$ to the eight rungs from $M=7$ to $M=25$ gives $\alpha = 1.406$ at $R^2 = 0.984$, and combined with Theorem 1 it predicts a scaling relation:

Proposition 16 (Error decays as a stretched exponential in area). *With $\varepsilon \propto 2^{-M}$ and $A \propto M^\alpha$, $\ln(1/\varepsilon) \propto A^{1/\alpha}$.*

Measured over $M \in [7, 25]$ — a sixfold range of area and a 250,000-fold range of error — the constant of proportionality holds to within 27% and drifts monotonically downward, so the relation is an approximation and we label it a proposition rather than a theorem. The drift has a cause, and finding it falsified a prediction we registered before synthesising.

The exponent is not constant, and the fit does not extrapolate. We predicted from the $[7, 25]$ fit that the reconciled TNF64, at $M = 56$, would cost 4,762 LUTs. It costs 7,479 at 48.20 MHz — the prediction is low by $1.57\times$. Reading α locally rather than globally shows why: it rises monotonically, from 1.06 between $M=7$ and $M=9$ to 1.78 between $M=15$ and $M=17$ and 1.85 between $M=25$ and $M=56$.

Theorem 17 (The ladder’s area law). *A rung’s area is a fixed overhead — registers, the exponent adder, the renormalisation — plus a significant multiplier quadratic in M :*

$$A(M) = f + cM^2.$$

Consequently the effective power-law exponent $\alpha(M) = d \ln A / d \ln M$ is not a constant but rises monotonically toward 2, and the marginal cost of a mantissa bit, $\lambda = dA/dM = 2cM$, is linear in M rather than fixed.

Synthesising the same width-derived multiplier at fourteen mantissa widths from $M=7$ to $M=90$ — 285, 372, 443, 567, 696, 870, 1238, 1683, 2601, 3990, 5934, 7479, 12232 and 19980 LUTs — and fitting both parameters gives $f = 141$ LUTs and $c = 2.4455$ LUTs per bit squared at $R^2 = 0.99963$, with a maximum deviation of 9.9% across a seventyfold range of area. Read locally, α climbs from 1.06 between $M=7$ and $M=9$ to 2.20 between $M=56$ and $M=70$, which is Theorem 17 rather than noise. On the registered prediction the structural model is right to +4.4% where the single power law was low by 36.3%; we report both, since the power law was ours and we published the prediction before synthesising.

Table 18: The same regime codecs valued at five mantissa widths through Theorem 18, with $c = 2.4455$ measured on XC7A200T. The unary regime’s silicon exceeds an entire TNF16 mantissa; the length-prefixed regime is below one bit everywhere above the 8-bit class.

M	λ (LUT/bit)	unary, 438 LUT	length-prefixed, 40 LUT
9	44.0	9.95	0.91
11	53.8	8.14	0.74
25	122.3	3.58	0.33
56	273.9	1.60	0.15
90	440.2	1.00	0.09

Theorem 18 (A regime’s worth in precision falls as $1/M$). *By Theorems 15 and 17, a regime codec of fixed area C is worth*

$$\frac{C}{\lambda} = \frac{C}{2cM}$$

mantissa bits, inversely proportional to the mantissa width. It is worth at least one mantissa bit exactly while $M \leq C/2c$.

Corollary 9 (Break-even widths). *With $c = 2.4455$, the unary regime is worth at least one mantissa bit of silicon up to $M = 89.6$ — that is, through the entire 128-bit class — while the length-prefixed regime reaches one bit only up to $M = 8.2$. takum’s regime therefore costs less than a single mantissa bit of silicon at every rung above TNF8, and posit’s costs more than one at every rung below TNF256.*

Corollary 9 is the exact form of what Corollary 11 below could only say qualitatively, and it is the sharpest statement we can make about where a fixed field pays for itself in area: against posit, almost everywhere; against takum, only at the very bottom of the ladder.

Corollary 10 (Range is paid for three times, and they are commensurable). *A tapered format pays for unbounded range in three currencies: word positions, in the regime field; silicon, in the codec; and precision at the extremes, in the taper. Theorem 15 and Theorem 1 put all three in mantissa bits, so they can be added.*

Applied honestly, Corollary 10 separates the two tapered families rather than lumping them. Against posit the silicon term alone is 8.98 bits at the 16-bit class — more than TNF16’s entire 9-bit mantissa, so the area a posit implementation spends decoding its regime would, spent on significand instead, roughly double the precision budget of a fixed-field format at the same class. Against takum the silicon term is 0.82 bits, which is small. The length-prefixed regime is cheap, and any argument for a fixed field made on area grounds is an argument against posit and not against takum. Against takum what remains is the accuracy result of Section 5, bounded as Theorem 12 requires to TNF’s own range.

Corollary 11 (The area argument is a narrow-width argument). *Since λ grows superlinearly in M while a regime codec’s cost is roughly fixed, the same codec converts to fewer mantissa bits as the class widens: 8.98 bits at 16 becomes 3.94 at 32 for the unary regime. The fixed-field area advantage is therefore strongest where the words are narrowest, which is the regime in which quantised inference operates and not the one in which numerical linear algebra does.*

14 The value law costs more than the field layout

Sections 12.1 and 13 priced the regime as a field to be extracted, and synthesising the published decoders shows that this is the smaller half of the question. We take the takum and posit

Table 19: Published 32-bit decoders, post-route on XC7A200T under one harness, no DSP, converting to binary32. These are the real codecs rather than the structural models of Table 16.

decoder	LUTs	RAMB36	$F_{\max}$	what dominates
posit32_decode	480	0	49.0 MHz	leading-zero count, shift
takum32_decode, published	76,821	0	—	2^f table, in logic
takum32_decode, clocked	10,967	84	11.4 MHz	same table, in BRAM
TNF32, exponent path	0	0	—	field read directly

decoders from an existing open verification suite — Verilog, golden-checked against an mpmath oracle — and put them through the same harness on XC7A200T.

The takum row is given twice because one number alone would mislead. Its decoder evaluates $2^{f_{hi}/2^{16}}$ from a 65536×48 -bit table, then applies a Taylor correction through a 108-bit and an 80-bit multiply. As published the table read is combinational, which no synthesiser can map to block memory, so it was built from logic at 76,821 LUTs. Registering that read and marking it `ram_style="block"` — one cycle of added latency, and a pipelined variant of the published module rather than the module itself — puts the table where it belongs: 10,967 LUTs and 84 RAMB36 tiles, which is 3.1 Mbit and 23% of this device’s entire block memory. We predicted 85 tiles from the table geometry before synthesising and measured 84. Neither number is a property of takum the format; both are properties of decoding a logarithmic value law into a linear datapath, and neither is the regime field.

Theorem 19 (Linear and logarithmic value laws). *Let a format’s value law be $v = s \cdot 2^e$ with s and e read from fields (linear), or $v = \pm 2^\ell$ for a decoded real ℓ (logarithmic). A linear law decodes with a shift, cost independent of the mantissa width. A logarithmic law requires evaluating an exponential, whose cost is a table or an iterative kernel and is independent of the field layout — therefore independent of the staircase form of Section 4.3.*

Theorem 19 is a second axis, orthogonal to the first. The staircase taxonomy classifies how a format spends *precision*; the value law classifies what it costs to *read*. posit is tapered and linear; takum is tapered and logarithmic; LNS is untapered and logarithmic; GF, TNF and IEEE are untapered and linear. The two axes are independent, and our earlier reading — that decode cost is set by whether the regime is scanned (Theorem 14) — describes only the field-extraction term. It remains correct within that term, and Table 19 shows the term is not where the money is for a logarithmic format.

Theorem 20 (The logarithmic decode costs by target precision, not by source width). *A 2^k -entry table for 2^f with a degree- d Taylor correction delivers*

$$P(k, d) = (d + 1) \left(k + \log_2 \frac{1}{\ln 2} \right) + \log_2 (d + 1)!$$

bits of relative accuracy, so the table is $P \cdot 2^k$ bits with $k = \frac{P - \log_2(d+1)!}{d+1} - \log_2 \frac{1}{\ln 2}$: exponential in the target precision, reducible only by raising the polynomial degree, which trades memory for multipliers. The width of the source format does not enter.

Proof. The residual after the table lookup is $x = f_{lo} \ln 2$ with $f_{lo} < 2^{-k}$, and the degree- d Taylor remainder of e^x is $x^{d+1}/(d+1)!$. Taking $-\log_2$ gives the stated P . $\square$ $\square$

Evaluated against direct simulation of the scheme — the table, the residual and the polynomial, in exact arithmetic — the formula holds to ± 0.1 bit at every point of a $k \in \{4, 8, 12, 16, 20\}$ by $d \in \{1, 2, 3, 4\}$ grid, up to the 51-bit resolution of the measurement itself.

Theorem 20 is falsifiable against the suite and survives. The takum32 decoder carries $d = 2$ and $k = 16$, and $48/3 = 16$ exactly, where 48 is the table’s output width. If the cost were set

by the source width, the takum64 decoder would need a larger table; if it is set by the target, it would need the same one. It reads `takum32_2frac.mem` — *the same file* — and emits binary32, as do the takum8, takum16 and takum64 decoders. One table serves four source widths because all four target one datapath.

Corollary 12 (The trade curve, and what each target actually costs). *Theorem 20 fixes a curve rather than a point, and the cheapest point on it that fits a given device is a computable quantity. Pricing the d multiplies at $2.4455P^2$ LUTs — the quadratic term of our own measured area law, Theorem 17 — and requiring the table to fit this device’s 365 RAMB36 tiles:*

<i>target</i>	<i>cheapest feasible point</i>	<i>block RAM</i>	<i>LUTs</i>
<i>binary32</i>	$d = 2, k = 15$	43 tiles	11,269
<i>binary64</i>	$d = 5, k = 16$	187 tiles	134,808

An XC7A200T has 134,600 LUTs, so decoding a logarithmic format into a binary64 datapath costs approximately the entire device.

We had previously written that binary64 decode “is not a design” at roughly 3.6 Tbit. That figure is the $d = 2$ corner of the curve and taking it for the whole curve was our error: at $d = 5$ the table is 6.9 Mbit and the design exists. The defensible statement is Corollary 12 — it costs a device — and not that it cannot be built.

The measured `takum32_decode` sits one k above the binary32 minimum, at $k = 16$ against 15, which is a $2\times$ memory margin and ordinary engineering practice rather than waste. A linear value law needs neither the table nor the polynomial at any target precision, because the significand is already in the datapath’s representation.

Corollary 13 (What TNF’s exponent path costs). *TNF’s value law is linear and its exponent field is fixed, so its exponent path is wires: no scan, no table, no iteration. Against posit32 that is 480 LUTs not spent; against takum32 it is a table of 3.15 Mbit not instantiated. This is a larger and more robust statement than the accuracy result, and unlike it, Theorem 12 places no range bound on it.*

We flag one asymmetry rather than let it pass. A logarithmic value law buys something real in exchange: multiplication becomes addition. The comparison above is of *decoders* — the cost of getting a value into a binary datapath — and a design that stays in the logarithmic domain never pays it. That is the regime in which LNS accelerators are built, and against those the argument of Corollary 13 does not apply at all.

15 Is there an optimal format, and in what sense

The results above are individually bounding: each says what a format cannot have, or what a claim may not be extended to. Read together they are stronger than that. They reduce the design space to four axes, three of which admit unconditional answers, so that once three inputs are named the format is determined rather than chosen.

Theorem 21 (The binary scale is the unique implementable optimum). *Let a format scale by r^e . Scaling is a shift — and by Theorem 19 free of any polynomial or table — if and only if $r = 2^k$ for integer $k \geq 1$. On that set $\kappa(r) = (r - 1)^2 / (r \ln r)$ is strictly increasing, taking 0.7213, 1.6230, 2.9455 and 5.0720 at $r = 2, 4, 8, 16$. Hence $r = 2$ minimises the expected relative error among all radices whose scaling is implementable without a transcendental, and it does so uniquely.*

The claim is verified by enumeration rather than resting on the monotonicity of κ alone. Sweeping r continuously over $[1.1, 16]$ with E integral and M taking the remaining width, at eight combinations of width and required range from $(P, R) = (8, 24)$ to $(32, 5925)$, the best implementable radix is $r = 2$ in every case, and the penalty against the unreachable optimum is $1.09\times$ at all eight — a constant, as it must be, since the two differ only by $\kappa(1.9)/\kappa(2)$. The unconstrained minimiser is not 2: it is near 1.9, worth 7.9% in error. That gain is unavailable, because any $r \neq 2^k$ makes r^e a transcendental evaluation, and Corollary 12 prices those at 84 RAMB36 tiles — 23% of this device — for a binary32 target. Trading 23% of a chip for 7.9% of a mantissa bit is not a trade anyone makes. Theorem 21 is therefore the interesting kind of optimality result: the optimum is not the extremum, it is the best point that can be built.

Corollary 14 (Three inputs determine the format). *The design space has four axes, and only one of them is a free choice.*

1. Radix of the scale. *Fixed at 2 by Theorem 21, unconditionally.*
2. Radix of the exponent’s encoding. *Fixed by the fabric. Theorem 8 shows a ternary exponent never spans more values per bit than a binary one when packed into bits, so it pays on a ternary fabric and nowhere else.*
3. Fixed fields or a taper. *A genuine choice, and by Theorem 6 exactly two options exist: bounded range at constant precision, or unbounded range at precision decaying no more gently than one bit per doubling (Theorem 12).*
4. The split between E and M . *Determined once a required range is named, by Theorem 1 — and only then, since the error depends on M alone and the range on E alone.*

Name the fabric, the required range, and which side of the dichotomy the workload sits on, and the format follows.

So the question “is there an ideal format” has an answer, provided it is asked with its three arguments supplied. Asked without them it is not underdetermined but ill-posed: Corollary 10 already says a budget cannot be recommended without a workload, and Theorem 6 says no single format can serve both sides.

Corollary 15 (What this makes TNF). *For the arguments (ternary fabric; a range the workload actually visits; fixed fields), Corollary 14 returns a binary-radix float with a balanced-ternary exponent field and a uniform mantissa split by $M = N - 1 - E_t$. That is TNF, and it is forced rather than designed. It follows that TNF is optimal exactly on those arguments and carries no claim off them — in particular none on a binary fabric, where axis 2 returns a binary exponent instead, and none against a tapered format outside TNF’s own range, where Theorem 12 forbids it.*

16 Where the format is the datapath

Every hardware figure so far compares a decoder against a datapath that contains a significant multiplier. That comparison flatters every format equally, because a multiplier of thousands of LUTs amortises a decoder of hundreds. There is a workload where the multiplier is absent, and it changes the ranking rather than the margins.

In a ternary-weight network — BitNet-style, with weights in $\{-1, 0, +1\}$ [30] — the product $w \cdot a$ is a select, a negate or a zero. No multiplier is instantiated, for any format. The datapath is *decode, apply the weight, add, accumulate*, and the decoder is now compared against an adder.

$2.03\times$ in area and $2.54\times$ in frequency, so $5.16\times$ in throughput per LUT — measured end to end rather than assembled from separate decoder and adder figures. The decomposition says

Table 20: One ternary neuron, matched width against matched width, every accumulator bit observed, median of five placement seeds, XC7A200T, no DSP. Every candidate decodes a *packed* word. Two corrections stand behind these numbers. An earlier version reported 440 against 895 LUTs, a $5.1\times$ ratio, which handed TNF pre-widened fields while the competitors unpacked theirs, and is withdrawn. The version after that reported $3.1\times$ at 16 bits and $5.6\times$ at 32 — on TNF modules that did not implement the specified format. At the specified parameters, and checked against a reference, the advantage is smaller and it is real.

width	format	LUTs	$F_{\max}$	MHz/LUT	ratio
8 bits	TNF8	545	62.13 MHz	0.1140	1.39 $\times$
8 bits	posit8	540	44.33 MHz	0.0821	—
16 bits	TNF16	565	66.28 MHz	0.1173	2.07 $\times$
16 bits	posit16	712	40.38 MHz	0.0567	—
32 bits	TNF32	569	66.91 MHz	0.1176	3.99 $\times$
32 bits	posit32	953	28.16 MHz	0.0295	—

why: against a 1,477-LUT multiplier and a 384-LUT adder, posit’s 456-LUT decoder is 20% of the datapath; with the multiplier gone it is 54%. For takum the same shift takes 85% to 97%, and brings 84 RAMB36 tiles with it.

Corollary 16 (Removing the multiplier promotes the decoder). *Let a datapath cost $C = D + M + A$ for decoder, multiplier and adder. A workload that eliminates M raises the decoder’s share from $D/(D+M+A)$ to $D/(D+A)$, and the ratio between two formats differing only in D from $1 + \Delta/(D_0+M+A)$ to $1 + \Delta/(D_0+A)$. The advantage of a decoder-free format therefore grows exactly as the multiplier shrinks, and is largest where there is none.*

Three things this does not say, and each would otherwise be found in minutes.

It is not that TNF avoids multiplication. It pays the significand multiply like every other format when one is needed; the quadratic term of its own area law (Theorem 17) *is* that multiplier. What a ternary network removes is removed for the network, not for the format.

It is not that TNF is better for FPGAs in general. We synthesised the exponent decode of TNF against an ordinary binary fixed field and they measure identically, 32 LUTs each, exactly as Theorem 8 requires. The advantage in Table 20 is over a *tapered* format, not over a fixed-field one: an IEEE-style field would place alongside TNF, not alongside posit.

And it is not a measurement on a ternary fabric. None exists to buy, and this is a binary FPGA where our own Theorem 8 says the ternary encoding earns nothing. It does not need to: the gain here comes from the decoder that is absent, and would be the same for any fixed-field format.

What the ternary fabric would add, structurally. On a fabric whose positions hold trits, a binary exponent field can be stored densely, in $\lceil b/\log_2 3 \rceil$ trits, or addably, one bit per trit — but not both. The dense form cannot be added without a radix conversion; the addable form wastes $1 - 1/\log_2 3 = 36.9\%$ of every position, which for the exponent widths in this catalogue is 33 to 57% of the field. A balanced-ternary exponent pays neither. This is structure rather than measurement, and it is the one place where the ternary encoding of Theorem 8 would stop being free of charge and start being free of cost.

Table 21: Dominance by width class. The family member that dominates is given for the hardest case in each class. Effective mantissa is measured; range is $3^{E_t} - 1$, which is analytic for a fixed field.

class	competitors	survivors	hardest case dominated by
8	7	0	$E_t=4, M=3$ against <code>fp8_e5m2</code>
16	7	0	$E_t=6, M=9$ against <code>takum16</code>
32	9	0	$E_t=5, M=26$ against <code>posit32</code>
64	4	0	$E_t=10, M=53$ against <code>cray_float</code>
128	1	0	$E_t=10, M=117$ against <code>binary128</code>

17 The ladder is a family, and the family is the frontier

The results above treat each rung as a format. Sweeping the exponent budget instead of fixing it changes what is being claimed, and strengthens it.

Theorem 22 (The width rule generates the frontier). *Fix a width N . The set $\mathcal{F}_N = \{(E_t, M) : 1 + E_t + M = N, E_t \geq 2, M \geq 1\}$ contains $N - 3$ members, each exact by construction. By Theorem 1 member E_t has effective mantissa $M = N - 1 - E_t$ and dynamic range $3^{E_t} - 1$ binades, so the members are totally ordered and no two dominate one another: $\mathcal{F}_N$ is a discrete Pareto frontier in $(M_{\text{eff}}, \text{range})$.*

The consequence is a claim about the catalogue rather than about one rung.

Corollary 17 (No catalogued competitor survives). *At every width class measured, each catalogued format of that width is strictly dominated by some member of $\mathcal{F}_N$: there is a member with at least its effective mantissa and at least its range, and strictly more of one.*

Measured against the oracles, twenty-eight competitors across five classes, with the dominating member named in each case:

Why this is not an artefact of density. A sufficiently dense family dominates any finite set of points trivially, so the content is in *why* the competitors sit below the frontier rather than on it. Each loses on at least one of three counts, and each count is one of the theorems above. A tapered format measures below its declared mantissa, by Theorem 12. A format that leaves positions unallocated pays 2^k for them, by Theorem 9 – which is what the historical GF-T ladder did at 2, 4 and 6 positions. And a binary exponent at equal position count spans fewer binades, since a trit carries $\log_2 3$ bits against a bit’s one. The family loses on none of the three because the width rule forbids the second and the encoding settles the third.

What this does not say. The pair $(M_{\text{eff}}, \text{range})$ is a property of the *number*, and dominance in it is not dominance in silicon. Section 16 measures the same formats through one datapath and puts TNF mid-pack among fixed fields at 0.114–0.118 MHz per LUT once each rung is measured at its specified parameters. Against `int8` the two are indistinguishable once seeds are swept (Table 43). The two are consistent: the frontier asks what a word of a given width carries, the datapath asks what it costs to unpack, and every fixed field unpacks by reading fields. The gap in silicon runs between fixed and tapered formats, 2.4 to 6.4×, and belongs to Theorems 14 and 19 rather than to this ladder.

A limit of the measurement, stated. Range is analytic for a fixed field, but verifying it at the boundary is not always possible: at $E_t = 49$ the edge of the range is $2^{1.4 \times 10^{23}}$, which no exact arithmetic can construct. Above roughly $E_t = 30$ the range is a fact about the encoding rather than a measured quantity, and is reported as such.

17.1 Which first principles

The width rule has been stated as a rule. It is not one: it is the solution of a constrained maximisation, and naming the constraint is what makes the rest of the paper follow rather than accumulate.

The rule is a Karush–Kuhn–Tucker condition. Write the design problem directly. Maximise precision subject to a position budget and a range that must be covered:

$$\max_{E,M} M \quad \text{s.t.} \quad 1 + E + M = N, \quad r^E \geq b + 1.$$

With multipliers λ on the budget and μ on the range, $\mathcal{L} = M + \lambda(N - 1 - E - M) + \mu(E - \log_r(b + 1))$ gives $\partial_M \mathcal{L} = 1 - \lambda = 0$ and $\partial_E \mathcal{L} = -\lambda + \mu = 0$, so $\lambda = \mu = 1$. Because $\mu > 0$, complementary slackness makes the range constraint *active*: at the optimum the exponent is exactly as wide as the range demands and not one position wider. That statement is Theorem 23, now derived rather than asserted. The multiplier is also the shadow price, and $\lambda = 1$ says a position of exponent costs exactly a position of mantissa — the exchange rate is unity because the budget is linear.

Why three and not two is a 1950s result, not ours. The radix economy $\rho(r) = r / \ln r$ is minimised at $r = e$. Ternary sits 0.46% above that optimum and binary 6.15% above it. We claim no credit for this; what the present work contributes is where it applies, which turns out to be the whole reconciliation of our two axes:

$$\underbrace{r}_{\text{cost per position}} \times \underbrace{\log_r V}_{\text{positions}}.$$

The number axis counts only the second factor, and there $\log_3$ against $\log_2$ gives the slack $0.3691 E$ of Corollary 26 unconditionally. The silicon axis restores the first factor, and packing a trit into binary fabric charges $\kappa(3)/\kappa(2) = 1.68$ per position, which is why BNF16 and TNF16 synthesise within 1% of one another. Theorem 8 and Corollary 26 are not in tension: they are the two factors of one product, reported separately.

What is derivation and what is only resemblance. Three connections are literal and are used as such: the multiplier argument above; radix economy, which fixes the exponent’s radix; and Kraft’s inequality, which bounds the taper in Theorem 12. Two resemblances are worth naming only so they are not mistaken for arguments. The variational form recalls least action, but there is no time and no trajectory here — the problem is a static optimisation, and the resemblance carries no consequence. And $1 + E + M = N$ is a budget, not a conservation law: positions are allocated by a designer, not withheld by nature. We state the distinction because a physical-sounding derivation would be doing rhetorical work that the multiplier argument already does honestly.

The standard falls out — at one choice of percentile, and not at others. Applied to the block axis the multiplier argument returns an element format, but which one depends on how much of the block distribution is protected, and that choice was left unstated in an earlier version of this section.

Protecting the median returns E1M2; protecting the 90th percentile or beyond returns E2M1, which is the element the OCP microscaling specification fixes. The correct statement is therefore a threshold rather than a derivation: *the width rule returns the MX element for any design protecting the 90th percentile of blocks or more, and a different element below that.* Which is right is a question about how many blocks may be allowed to saturate, not one settled from first principles, and we say so rather than quoting the agreeing case alone.

Table 22: The width rule on the block axis. The measured within-block span sets $E^* = \lceil \log_2 b \rceil$, and the answer moves with the percentile protected.

percentile protected	span, binades	E^*	element
50th (median)	1.89	1	E1M2
90th	2.45	2	E2M1
99th	3.04	2	E2M1
99.9th	3.75	2	E2M1

Applied to the block axis this stops being abstract. The within-block span of a trained network’s weights, measured before any format is chosen, is 1.89 binades at the median and 3.04 at the 99th percentile. The optimum is $E^* = \lceil \log_2 b \rceil$, so protecting the median returns E1M2 and protecting the 99th percentile returns E2M1 — which is the element the OCP microscaling specification fixes. The standard is not reproduced by coincidence: it is the unique solution of the multiplier problem at the tail these weights actually have. Theorem 24 says which of the two should win, since under-sizing pays on the tail linearly while over-sizing pays logarithmically on every value, and Section 23.7 reports the measurement that decides it.

17.2 The width rule is an optimisation with one solution

Theorem 46 compares against a given competitor. Stated as a design problem it selects a member outright, which is what makes the rule falsifiable: it names a winner before any measurement.

Theorem 23 (Optimal member for a named range). *Let a workload visit b binades. Among all members of $\mathcal{F}_N$ that represent it without overflow, the effective mantissa is maximised uniquely at*

$$E_t^* = \lceil \log_3(b+1) \rceil, \quad M^* = N - 1 - \lceil \log_3(b+1) \rceil.$$

Proof. Representability requires $3^{E_t} - 1 \geq b$, i.e. $E_t \geq \log_3(b+1)$, and E_t is an integer, so $E_t \geq E_t^*$. By the width rule $M = N - 1 - E_t$ is strictly decreasing in E_t , so the constraint binds and the minimiser is unique. $\square$

The rule therefore has no free parameter once the workload is named, and naming the workload is a measurement rather than a choice. Section 16 uses this as a prediction: the binade span of a network’s weights is measured first, E_t^* is computed from it, and only then is perplexity evaluated across the whole family. A different winner would falsify the rule.

Theorem 24 (Regret of a misnamed range). *Sizing for b' when the workload visits b costs*

$$R = \begin{cases} \lceil \log_3(b+1) \rceil - \lceil \log_3(b'+1) \rceil & \text{mantissa bits, } b' > b, \\ \text{overflow on a set of measure } 1 - \frac{\log_3(b'+1)}{\log_3(b+1)}, & b' < b. \end{cases}$$

The penalty is asymmetric: over-sizing costs precision logarithmically, under-sizing costs range linearly, so when the visited range is uncertain the rule should round up.

The asymmetry is worth naming because it contradicts the intuition that a wider exponent is the conservative choice. It *is* conservative — but the price is paid in every value the format ever represents, whereas the price of being too narrow is paid only on the tail. Which dominates is an empirical question about the tail’s weight, not a matter of taste.

Corollary 18 (When the trit is worth its packing). *Ternary buys $\Delta = E(1 - \log_3 2)$ positions and costs $\kappa(3)/\kappa(2) = 1.683$ per position when packed into binary fabric, where $\kappa(r) = (r -$*

Table 23: The four families at eight bits, one workload, one metric. SmolLM2-135M on wikitext-2, 40 windows of 2048 tokens, per-tensor scale, baseline 14.4874 verified in band before any comparison. The predictions were printed *before* the perplexities were computed.

family	rule, radix	magnitudes	ppl	vs fp32
fp32 reference	—	—	14.4874	1.000×
GF8 , $E=3$	golden ratio, binary	129	14.6130	1.009×
GF-T8, $E_t=3$	golden ratio, ternary	109	15.5147	1.071×
BNF8 , $E=3$	width rule, binary	129	14.6130	1.009×
TNF8 , $E_t=2$	width rule, ternary	73	14.7012	1.015×

Table 24: The GF-T and TNF ladders rung by rung. Eight of the nine specified rungs are identical — the two rules choose the same budget — and they part only at 32 positions.

rung	GF-T (E_t, M)		TNF (E_t, M)		
4	2	1	2	1	identical
8	3	4	3	4	identical
16	4	9	4	9	identical
32	5	21	6	25	the one divergence
64	7	52	7	52	identical
128	8	115	8	115	identical
256	9	242	9	242	identical
512	10	497	10	497	identical
1024	11	1006	11	1006	identical

$1)^2/(r \ln r)$. The two cancel at E satisfying $\Delta = 0.331 E_t$, so on positions the trit wins outright and on binary bits it breaks even — which is why BNF16 and TNF16 synthesise within 1% of one another while the number-axis frontier separates them.

18 Four families on two axes

TNF is one corner of a square, and the square is worth drawing because the two rules that define its axes were derived separately and for unrelated reasons.

One axis is the *rule that fixes the exponent budget*. The golden-ratio rule takes $E = \text{round}((N - 1)/\varphi^2)$ and produces the GF ladder; the width rule takes $1 + E + M = N$ and produces the BNF ladder. The other axis is the *radix of the exponent’s encoding*: binary, or balanced ternary. Crossing them gives four families — GF and GF-T on the golden-ratio rule, BNF and TNF on the width rule — and each is a distinct format rather than a renaming of another.

Two names, one encoding, two rules. The square’s corners are not four unrelated formats. GF and BNF share the binary encoding and differ in which rule fixes E ; GF-T and TNF share the ternary encoding and differ in the same way. How far the rules then diverge is a measurement, and mostly they do not.

This is why the silicon tables carry TNF rows and no separate GF-T ones: at every rung but 32 they would be the same numbers twice. It is also why the eight-bit measurement below reports GF-T8 and “TNF8 at $E_t = 3$ ” as one line of 109 magnitudes and 15.5147 — they are the same format, reached from two directions. BNF is measured too, at 97 LUT and 388.35 MHz,

but on the isolated-decoder harness of Section 7 rather than the ternary-neuron datapath of Table 39, which is where it now appears in Table 39 as row 2, measured on the same harness as everything else.

The two rules name the same format, and neither predicted it. GF8 at $E = 3$ and BNF8 at $E = 3$ return 14.6130 — the same number, not a close one, because at eight bits the two derivations select the identical layout, $E = 3$ and $M = 4$. That is a stronger statement than either rule alone and one we did not set out to find.

And the width rule’s prediction was falsified twice. It named BNF8 $E = 4$ and TNF8 $E_t = 3$; the measured winners are $E = 3$ (14.6130 against 14.6592) and $E_t = 2$ (14.7012 against 15.5147). Both predictions were one step too wide. The rule’s *form* survives and is visible in the sweep — a single optimum, both directions from it worse, and the penalty asymmetric exactly as Theorem 24 requires, with $E = 1$ giving 4.5 million and $E = 2$ giving 288 while $E = 4$ costs 0.3% and $E = 5$ costs 6%. What was wrong is the estimator of the range the workload visits, not the rule that consumes it: we measured the span from the 0.1st percentile to the maximum, which credits a tail carrying almost no energy. Recorded rather than quietly re-tuned, because a rule that names its winner before the measurement can be wrong, and this one was.

Ternary loses on a binary fabric here too, for the third independent time. GF-T8 carries 109 magnitudes where GF8 carries 129, and pays 15.51 against 14.61; TNF8 at $E_t = 2$ carries 73. Three measurements now agree from different directions: BNF16 and TNF16 within 1% in placed silicon, GF8 against GF-T8 here, and MXFP4 against TNF4 on the block axis (Section 21). The ternary exponent earns on positions and pays on bits, exactly as Theorem 8 requires, and this fabric charges in bits.

19 Two formats close a ternary node

A ternary node contains exactly three sites where a number format could live, and only one of them needs one.

1. The **weight** is a code. Multiplying by it is a sign-select $(+x, -x, 0)$, not a multiply.
2. The **sample** arrives ADC-native and is used as it lands.
3. The **accumulator** is the only object in the system carrying a dynamic range that must be spent.

This is the whole design problem, and it is smaller than the literature makes it look. The published ternary methods answer (1) and leave (3) in `fp16` or `int8`; the published format work answers (3) and assumes a multiplier at (1). We answer both, with two formats that were derived independently and turn out to be complements.

19.1 Why the weight alphabet is φ , and why that is forced

GFternary is a two-bit alphabet $\{-\varphi, 0, +\varphi\}$. The φ is not decoration, and the following says why it is the only possible choice.

Theorem 25 (The golden alphabet is unique). *Let a three-symbol weight alphabet be $\{-r, 0, +r\}$ with $r > 1$, and require that the product of two weights be expressible in the additive lattice the datapath already computes in, i.e. $r^2 = r + 1$. Then $r = \varphi$ uniquely.*

Proof. $r^2 - r - 1 = 0$ has roots $(1 \pm \sqrt{5})/2$, of which one is positive. □

Corollary 19 (Depth costs no multiplier). *The gain of a stack of k ternary φ -layers is exactly*

$$\varphi^k = F_k \varphi + F_{k-1},$$

a pair of integers with F the Fibonacci sequence. Rescaling between layers is therefore a shift-and-add, and the zero-multiplier property survives to arbitrary depth.

This is the practical content of Theorem 25, and it is what separates the φ alphabet from the $\{-1, 0, +1\}$ one used throughout the ternary-network literature. With unit weights the layer gain is 1 and carries no information, so every published method attaches a learned real scale α_ℓ per layer — and multiplying by α_ℓ puts the multiplier back. With the φ alphabet the inter-layer scale is known exactly and is a pair of integers, so it is never multiplied by. Storage is two bits either way and the symbol count is three either way; the φ alphabet simply carries a scale that the unit alphabet has to learn and then pay for.

19.2 The linear path is exact

Theorem 25 is usually read as removing a multiplier. It removes the rounding with it: on the alphabet it fixes, the entire linear path of a network is computed exactly.

Represent a value as an integer pair (a, b) standing for $a + b\varphi$. Then

$$\varphi(a + b\varphi) = a\varphi + b\varphi^2 = a\varphi + b(\varphi + 1) = b + (a + b)\varphi,$$

so applying a weight is the map $(a, b) \mapsto (b, a + b)$ — the Fibonacci recurrence, one integer addition, no shift. Negation flips both components and a zero weight is a skip.

Theorem 26 (Exactness of the multiply-free path). *$\mathbb{Z}[\varphi] = \{a + b\varphi : a, b \in \mathbb{Z}\}$ is a ring, and the weight alphabet $\{-\varphi, 0, +\varphi\}$ lies inside it. Hence for inputs in $\mathbb{Z}[\varphi]$ the entire linear part of a ternary network — every weight application and every accumulation, to arbitrary fan-in and arbitrary depth — remains in $\mathbb{Z}[\varphi]$ and is computed without rounding error of any kind.*

Proof. Closure under addition is componentwise. Closure under multiplication follows from $\varphi^2 = \varphi + 1$, which reduces every product to the stated integer-pair map. The accumulator is therefore a pair of exact integer registers. $\square$

Measured on a layer of fan-in 512 with weights drawn from the alphabet, the integer pair reproduces the real-valued sum to the precision of the checking arithmetic and not to the precision of the datapath: there is nothing in the datapath to be imprecise. The component widths grow logarithmically, reaching eight bits over those 512 terms.

19.3 Closure is what removes the machinery

The multiply-free datapath is not new as an aim, and the published attempts make the role of Theorem 26 visible by contrast.

The closest is FQP [28], a Fibonacci quantization processor that turns the products of Fibonacci-quantized weights into additions. To do so it carries two distinct arithmetic units — a dualistic-transformation adder for the large operands and a bit-exclusive adder for the small ones — and a topological-order routing scheme that maps data onto them. Fibbinary quantization for neural radio receivers [29] takes the softer route on the same substrate and reports 45% and 44% savings in multiplier power and area: the multiplier is made cheaper, not removed.

The machinery in the first is not an artefact of implementation. It follows from a property of the number set.

Corollary 20 (Closure removes the transformation stage). *Let A be a weight alphabet and let s be the operation that applies a weight. If the additive group generated by A is closed under s , the datapath needs only that group’s addition. If it is not, every product must be re-expressed in the representable set, which requires a transformation stage and a schedule that routes operands to it.*

Proof. Immediate from Theorem 26 and its negation. For $A = \{-\varphi, 0, +\varphi\}$ the generated group is $\mathbb{Z}[\varphi]$ and $s = \varphi \cdot$ maps $(a, b) \mapsto (b, a + b)$, which lands inside it, so addition suffices. For A a set of Fibonacci integers $\{F_k\}$, the product $F_i F_j$ is in general not a Fibonacci integer — $F_4 F_4 = 9$ is not in $\{1, 1, 2, 3, 5, 8, 13, \dots\}$ — so the product leaves the representable set and must be transformed back into it. $\square$

19.4 How many multiply-free scales are there

Closure explains why φ works. It does not by itself say φ is the only choice, and the honest answer is that it is not — it is the only choice at its size.

A scale r can be applied without a multiplier exactly when r is an algebraic integer whose companion matrix has entries in $\{0, \pm 1\}$: writing $r^d = a_{d-1}r^{d-1} + \dots + a_0$, multiplication by r maps the coordinate vector $(x_0, \dots, x_{d-1})$ to $(a_0 x_{d-1}, x_0 + a_1 x_{d-1}, \dots)$, and any $|a_i| > 1$ reintroduces a constant multiply. The condition is finite, so the scales can be enumerated rather than argued about.

Theorem 27 (The multiply-free scales, by degree). *Over coefficients in $\{0, \pm 1\}$:*

- **Degree 1** admits $r \in \{0, \pm 1\}$, which is no scaling, and $r = \pm 2^k$, free by shifting. One register, ratio 2.
- **Degree 2** admits exactly two equations with a real root of modulus other than 0 or 1: $r^2 = r + 1$ and $r^2 = -r + 1$, whose roots are φ and φ^{-1} — the same ladder. One addition, two registers, ratio 1.618.
- **Degree 3** admits three: $r^3 = r + 1$ (plastic, 1.3247), $r^3 = r^2 + 1$ (supergolden, 1.4656), $r^3 = r^2 + r + 1$ (tribonacci, 1.8393). At most two additions, three registers.

Proof. Exhaustive over the nine degree-2 and twenty-seven degree-3 coefficient triples; the degree-2 discriminants $a^2 + 4b$ are negative for four cases, and the remaining five give roots in $\{0, \pm 1\}$ except for $b = 1, a = \pm 1$. $\square$

Corollary 21. *Nothing lies between the shift and φ . φ is the unique multiply-free refinement of the power-of-two ladder available at two registers.*

So the multiply-free scale is not a point but a hierarchy indexed by algebraic degree, and each degree buys finer granularity for one more register. Table 25 measures the trade at the first two rungs where it is a trade at all. The plastic step $(a, b, c) \mapsto (c, a + c, b)$ costs the same single addition as the Fibonacci step and one more register, and its ladder is finer: successive levels differ by 13.97% against φ ’s 23.61%.

At 32 bits the refinement costs 2.2% of the area and 6.4% of the frequency. That is close enough to free that the choice is a question about the network rather than about the arithmetic, so we asked the network.

20 Which rung to use, and how the answer moves

Table 26 quantises every 2-D linear weight of SmolLM2-135M to $\pm r^k$ with a per-output-channel scale, the codebook sized to exactly 2^b entries including zero and sign, and evaluates wikitext-2 perplexity over twelve windows of 2048 tokens. Baseline fp32 is 14.36.

Table 25: One scale application, xc7a200t, median of five seeds. Matched: neither unit has a direction mux. Level error is $(r - 1)/(r + 1)$, the worst relative gap on a ladder of ratio r .

scale	LUT	FF	$F_{\max}$ (MHz)	level error
φ , 16-bit	149	128	307.69	23.61%
plastic, 16-bit	186	160	318.47	13.97%
φ , 32-bit	223	192	247.10	23.61%
plastic, 32-bit	228	200	231.21	13.97%

Table 26: Perplexity by ladder and code budget, SmolLM2-135M on wikitext-2. Coarse ladders span range and waste precision; fine ones do the reverse. The best rung moves down the hierarchy as the budget grows.

bits	shift (2)	φ (1.618)	supergold (1.4656)	plastic (1.3247)
3	2309.86	41835.53	1367268.11	6720092.29
4	76.81	24.43	25.10	55.72
5	77.52	22.73	18.93	16.45

The same sweep on Qwen2.5-0.5B — a different architecture at four times the parameters, with a different tokenizer and different training data — returns the same winner at every budget and the same complete ordering at every budget: shift 2384.96, then φ **16.24** against supergolden 17.98 and plastic 47.41, then plastic **13.17** against supergolden 14.09 and φ 15.33, from an fp32 baseline of 12.27. At four bits the larger model separates φ from supergolden by 10.7%, where the smaller one gave 2.7% and no separation.

A ladder of ratio r with n magnitudes spans r^{n-1} , so finer steps are bought with span. At three bits nothing spans enough and the coarsest wins; at four the balance sits at φ ; at five the finest tested wins outright. This is the range-against-precision law of Section 4 applied to the scale hierarchy rather than to a field layout, and it says plainly that φ is optimal at a budget, not everywhere. At four bits the two models bracket the comparison: SmolLM2 separates φ from supergolden by 2.7%, which twelve windows cannot resolve, while Qwen separates them by 10.7%, which they can. φ is at worst tied there and ahead on the larger model, and supergolden costs a register either way. One bit up, the plastic number wins on both, by 28% and 14%, with φ third of four — and by then the affordable part of the hierarchy has been passed as well.

The choice does not need the network. Ranking the ladders by mean squared quantisation error on the weights alone reproduces the perplexity order exactly at three and five bits, and at four bits transposes only the pair perplexity cannot separate.

Theorem 28 (The ladder law). *Among multiply-free ladders at a fixed code budget, the one minimising weight mean squared error minimises perplexity, up to pairs whose error differs by less than the evaluation resolves.*

The error itself has a closed form, so the rung follows from a histogram. Normalising each channel by its maximum puts the levels at r^{-k} for $k = 0, \dots, n - 1$ with $n = (2^b - 1)/2$, and

$$\text{MSE}(r, b) \approx c(r)^2 \mathbb{E}[x^2 \mathbf{1}_{x>t}] + \mathbb{E}[x^2 \mathbf{1}_{x<t}], \quad t = \frac{1}{2} r^{-(n-1)},$$

where $c(r)^2$ is the mean square *relative* rounding error of a geometric ladder of ratio r — a one-dimensional integral over log-position within a bin, independent of the data. The first term falls as r falls and the second rises as the clipping threshold does; the minimum is the crossover. Against the measured error over 106,168,320 normalised weights the form agrees to

within 1–9% at four and five bits and selects the same ladder at all three budgets. At three bits it underestimates by a third, for a reason worth stating: at seven codes the regime is absolute truncation, not relative rounding, and the model is built on the latter.

20.1 What the four-bit transposition is, and what it is not

The form above ranks correctly at three and five bits and transposes the top pair at four, choosing supergolden where perplexity chooses φ . The transposition reproduces on both models in the same direction, so it is an effect. Two candidate mechanisms were tested; one failed and one held, and the pair is more informative than either alone.

Activation salience: fails. The output error of a layer is $\sum_j \delta w_{ij} x_j$, so a weight multiplying a large input channel should cost more, and the channels are far from equal — the first attention projection spans $28\times$ between its largest input channel and its median on SmolLM2-135M, $22.4\times$ on Qwen2.5-0.5B. This is the observation AWQ [27] is built on. Ranking by $\sum_j a_j^2 \delta w_{ij}^2$ reverses the four-bit pair on SmolLM2, matching perplexity — and does not on Qwen, where the weighting scales all four candidates by very nearly 1.09 and therefore cannot reorder them. An explanation that needs channel salience to correlate with the error pattern is not general, because in one of two models it does not.

Loss curvature: holds. The quantity the transposition is about is not how large an operand is but how much the loss moves when a weight moves. Ranking by the diagonal Fisher weighting $\sum_{ij} g_{ij}^2 \delta w_{ij}^2$ with $g = \partial L / \partial w$ reverses the four-bit pair to φ on *both* models and leaves five bits at plastic on both, matching perplexity in all four cases.

Theorem 29 (Curvature correction). *Among multiply-free ladders at a fixed code budget, the minimiser of $\sum_{ij} g_{ij}^2 \delta w_{ij}^2$ predicts the perplexity ranking where the unweighted error does not, and does so on both models tested. The correction moves the choice toward the coarser ladder.*

The two results together say something the successful one alone would not. Activation magnitude is half of the gradient and not the half that decides: $g_{ij} = \delta_i x_j$ carries both the input the weight multiplies and the sensitivity of the loss to that unit’s output, and it is the second factor that makes the correction replicate. A method calibrated on activations alone — and several published ones are — is using a proxy that agreed with the target on one of our two models.

Theorem 30 (Optimal ratio). *For a weight distribution and a code budget, the optimal multiply-free ratio is the minimiser of $\text{MSE}(r, b)$, and the degree required is the smallest whose minimal scale does not exceed it. Minimising over continuous r gives $r^* = 2.054, 1.4115, 1.1999, 1.1050, 1.0576$ and 1.0319 at 3 to 8 bits, with $\ln r^*$ halving per bit as doubling the code count halves the available log-spacing.*

21 The block axis, and the half of it that is ours

Everything above concerns per-channel scaling. The other deployed axis shares one scale across a short block, and there our formats lose. It is worth stating how thoroughly, because the answer is not a ranking but a ceiling.

Lloyd–Max on the within-block distribution — 18,850,950 values, block of 32 along the contraction axis, E8M0 shared scale as the MX specification mandates — gives the eight-level codebook minimising squared error. It is neither multiply-free nor a ladder; it is the best any eight magnitudes can do.

Table 27: Block axis, SmolLM2-135M, 40 windows, fp32 baseline 14.4874. The optimal codebook improves squared error nearly fifteenfold and perplexity by 0.9%, and every ladder that beats MXFP4 on the first loses on the second.

element, 8 magnitudes	squared error	$\times$ bound	ppl
Lloyd–Max (optimum, not implementable)	1.341×10^{-3}	1.00	22.2976
MXFP4 (E2M1)	2.007×10^{-2}	14.97	22.4998
supergolden/plastic, 4 fine steps	2.068×10^{-3}	1.54	27.3079
supergolden/plastic, 3 fine steps	2.116×10^{-3}	1.58	30.8306
φ /supergolden, 2 fine steps	2.749×10^{-3}	2.05	38.6471
φ /plastic, 2 fine steps	2.823×10^{-3}	2.11	55.4495

Theorem 31 (The block axis is not an element-format problem). *On a short block with a shared scale, the element codebook minimising squared quantisation error is not the one minimising perplexity, and the gap between the best achievable codebook and a deployed one is under one percent. Effort spent designing element formats for this axis cannot recover more than that.*

The ranking does not merely fail to transfer; it inverts. Every two-segment ladder here beats MXFP4 on squared error by factors of three to ten, and every one loses on perplexity by 21% to 146%.

The reason is what a block is. Thirty-two weights share one scale, and squared error charges a clipped small weight at its own magnitude — so a codebook concentrating its levels near the top of the range scores well, having put its resolution where the mass is. But the network does not read a block’s weights independently; they are summed against activations, and a codebook that resolves the top of every block finely while collapsing its lower two-thirds destroys the block’s ability to represent a *direction*, not merely its magnitudes. E2M1 spends 0, 0.5, 1, 1.5, 2, 3, 4, 6: a $12\times$ span with levels at both ends. It scores badly on squared error precisely because it refuses to concentrate, and it wins on perplexity for the same reason.

We record this as the competition’s result, not ours. Three of our own attacks failed here — a ternary-exponent float, constant-ratio ladders, and two-segment ladders built to imitate E2M1’s spacing — and the bound explains why a fourth is not warranted.

21.1 The other half of the block axis

The bound closes the element format. It says nothing about the scale, which is one value per thirty-two weights and had never been varied. Its cost is amortised thirty-twofold, which is the regime where a finer grid is affordable — and the accounting decides the comparison before any perplexity is measured.

E8M0 gives the scale eight exponent bits. Over 3,317,760 block scales the range actually occupied is 8.32 binades on SmolLM2-135M and 9.12 on Qwen2.5-0.5B, so four bits suffice. At four bits 2^k needs 8.3 steps of the sixteen available and φ^k needs 12.0, with half the step. The choice at equal width is between exactly those two.

A stronger reading of the first column does not survive the second. On SmolLM2 the response is non-monotone with its minimum at φ and $2^{k/4}$ falling below plain 2^k , which would say φ is the optimum of the grid; on Qwen the response is monotone in granularity and the finest grid wins. The sign replicates and the size does not — 5.09% against 0.63% — and the claim of optimality does not replicate at all. We state only what does.

The frontier, not the pair. A scale is one value per K weights, so it costs bits/ K per weight, and comparing formats in pairs hides that. Table 29 puts perplexity against total bits per weight. The Pareto set is identical on both models.

Table 28: Block scale grid, K=32, E2M1 elements, 40 windows. At equal bits φ^k beats 2^k on both models; finer grids cost one or two more bits and their benefit does not replicate.

scale grid	bits	SmolLM2	Qwen
2^k (the MX specification’s)	4	22.4998	14.9447
φ^k	4	21.3545	14.8512
$\sqrt{2}^k$	5	21.8960	14.7456
plastic ^k	5	22.7625	15.0333
$2^{k/4}$	6	24.0791	14.4328

Table 29: Scale-cost frontier, E2M1 elements, 40 windows. φ^k at four bits per block of 32 strictly dominates MXFP4 — cheaper and better on both models. Above 4.25 bits per weight the frontier belongs to scales carrying a mantissa.

scheme	scale b/w	total b/w	SmolLM2	Qwen
φ^k 4b/64	0.0625	4.0625	22.0978	15.0090
φ^k 4b/32	0.1250	4.1250	21.3545	14.8512
2^k 4b/32	0.1250	4.1250	22.4998	14.9447
MXFP4 , E8M0 8b/32	0.2500	4.2500	22.4998	14.9447
E4M3 8b/32	0.2500	4.2500	19.8628	13.7636
NVFP4-like E4M3 8b/16	0.5000	4.5000	18.5445	13.5340

Three things follow, and only the third is ours.

MXFP4’s scale field is wider than it needs to be. That E8M0’s range exceeds what weights occupy has been observed elsewhere [25] and we claim no priority for the observation; what is measured here is how much, and what the freed bits are worth. 2^k at four bits reproduces E8M0’s perplexity *exactly* — 22.4998 and 14.9447, to the last digit — because the occupied range is 8.32 and 9.12 binades and eight exponent bits buy nothing over four.

At those four bits φ^k beats 2^k , so φ^k 4b/32 strictly dominates MXFP4: cheaper by 0.125 bits per weight and better on both models. It is not a trade.

And the upper frontier is reached by spending more. E4M3 — a scale carrying a mantissa, which is what NVFP4 [26] uses — reaches 19.8628 and 13.7636 at the same 4.25 bits per weight where MXFP4 sits, against our 21.3545 and 14.8512 at 4.1250. The comparison is not at equal width, and the next subsection shows what happens when it is made so.

21.2 A geometric scale grid dominates a float one

A scale is read as a multiplier, so the error that matters is relative, and the grid minimising relative error over a log-distributed range is uniform in log. A float scale is not.

Theorem 32 (The geometric scale grid). *A float scale $2^e(1 + m/2^k)$ places 2^k points in each binade, with largest adjacent ratio $1 + 2^{-k}$ at the bottom of the binade. A geometric grid with the same point count places them at $2^{j/2^k}$, with every adjacent ratio equal to $2^{2^{-k}}$. Since $2^x = e^{x \ln 2} < 1 + x$ for $x > 0$, the geometric grid’s worst-case step is strictly smaller at every mantissa width, and the ratio of the two excesses rises monotonically to $1/\ln 2 = 1.442695$.*

Proof. Immediate from $\ln 2 < 1$. The limit follows from $(1 + x) - 1/(2^x - 1) \rightarrow 1/\ln 2$ as $x \rightarrow 0^+$. $\square$

The defect is structural rather than a matter of tuning: a float grid is uniform in *value* within a binade and therefore non-uniform in log, spending resolution at the top of each binade

Table 30: Scale grid at equal width, K=32, E2M1 elements, 40 windows. A geometric grid beats E4M3 at both widths on both models, as Theorem 32 requires. The margin does not transfer between models; the sign does.

scale grid	bits	SmolLM2	Qwen
geometric	7	18.8024	13.6401
E4M3	7	19.8628	13.7636
geometric	8	18.1238	13.6910
E4M3	8	19.8628	13.7636

where the relative step is already small. Root-mean-square relative error over a binade agrees — 0.1138 against 0.1068 at two mantissa bits, 0.0550 against 0.0516 at three, 0.0270 against 0.0254 at four.

Measured at equal width, on both models, the prediction holds (Table 30).

One further reading is model-dependent and we do not make it. On SmolLM2 a geometric grid at 4.25 bits per weight beats the NVFP4-like configuration at 4.50 — cheaper and better; on Qwen it is cheaper and worse, 13.6910 against 13.5340. *Geometric beats float at equal bits* replicates; *geometric at lower cost dominates NVFP4’s configuration* does not.

The tension this exposes. The optimal grid is geometric with ratio $R^{1/N}$ for the occupied range R and point count N , and that ratio is an arbitrary real: applying it needs a multiplier. φ^k is geometric *and* multiply-free, at a ratio fixed at 1.618 which is best where the range and the count want it — at four bits, which is where it wins above. So the placement is a pair of statements rather than one. Among all scale grids, geometric beats float and the best geometric grid needs a multiplier; among multiply-free grids, φ^k is the finest available at four bits. The first is a contribution to the field, the second is what this datapath gets for nothing.

This is not our format beating MXFP4: the element stays E2M1, and only the base of the shared exponent changes, so it is a contribution to the microscaling family rather than a defeat of it. And the grid that helps is ours specifically — on a binary datapath 2^k is a shift and φ^k is not, while on the datapath of Section 19.3 φ^k is k Fibonacci steps and needs no multiplier. A scale grid that costs everyone else a multiplier is free here.

Scope. This is stated for per-channel scaling and does not extend to short blocks. Applied to the within-block distribution under a block of 32 and an E8M0 shared scale, the same form returns $r^* = 1.3308$, which measures 37.06 against the 23.36 of the ladder it ranks worse — and perplexity there is not even monotone in r , running 39.09, 23.36, 39.61 as the ratio falls through φ , supergolden and plastic. The form charges clipping at the weight’s own squared magnitude, which is right when a scale is shared by thousands of weights and the clipped ones are a tail, and wrong when thirty-two share one and the clipped ones are a fifth of the block.

22 Fineness costs registers, not adders

Reading the hierarchy by degree, and taking the cost to be the minimal ratio available at each degree, hides the result. Read along a family instead.

What is forced, and what is not. For base φ , three digits is the proven minimum for constant-time carry-free addition [21, 19]; a two-digit alphabet is impossible. $\{-1, 0, 1\}$ is *one* of the minimal alphabets — chosen because it is symmetric, so subtraction reuses the adder — and not the unique one: $\{0, 1, 2\}$ is minimal too, and a working $\{0, 1, 2\}$ parallel adder for

Table 31: The multiply-free scales in one place. Level error is $(r - 1)/(r + 1)$, the worst relative gap on a ladder of ratio r ; silicon is one scale application at 16-bit components, xc7a200t, median of five seeds. Every rung below the shift costs exactly one addition; only the register count grows.

scale	polynomial	r	regs	adds	level err.	LUT / $F_{\max}$
shift	$r = 2$	2.0000	1	0	33.33%	—
φ	$r^2 = r + 1$	1.6180	2	1	23.61%	149 / 307.69
supergolden	$r^3 = r^2 + 1$	1.4656	3	1	18.88%	—
plastic	$r^3 = r + 1$	1.3247	3	1	13.97%	228 / 231.21
	$r^5 = r^3 + 1$	1.2365	5	1	10.57%	164 / 306.56
	$r^4 = -r^3 + r^2 + r + 1$	1.1787	4	3	8.20%	469 / 184.98
	$r^6 = r + 1$	1.1347	6	1	6.31%	—
	$r^8 = r + 1$	1.0970	8	1	4.63%	225 / 815.66

φ was implemented and tested here. Légerský [20] extends the bound to any finite alphabet inside $\mathbb{Z}[\varphi]$ containing zero, so no non-integer set escapes it.

The bound constrains the *digit* alphabet of a positional representation and one operation — rewriting a digit-wise sum back into the alphabet with a sliding window. A weight set has no base and no carry, so it does not transfer to the weight alphabet of a dot product. It would transfer only if the accumulator were itself a redundant base- φ positional format, and ours is a pair of integers never reduced to a digit string. We record the theorem because it says something real and narrow — two letters cannot be made to work at any window size — and record its scope because the wider reading, which we first wrote down, is false.

The window is not small either: the published algorithm on $\{-1, 0, 1\}$ is 21-local, with memory 10 and anticipation 10. “Constant time” read as “cheap” would be wrong by an order of magnitude.

A qualification the table needs. Not every rung is Perron — that is, not every one has its real root strictly dominant among the conjugates. The degree-4 minimum has a conjugate of modulus 1.5129 against a root of 1.1787, and the degree-6 minimum 1.3069 against 1.0863; the degree-5 minimum 1.1237 is Perron, its largest conjugate being 1.0959. Where the root is not dominant, the coordinate vector of r^k grows as the larger conjugate rather than as the value, so a chain of k applications needs $k \log_2(\rho/r)$ guard bits — 0.36 bits per step at degree 4. That is a cost on composition and not on a single application, which is why the silicon column does not show it, and it is stated here rather than left for a reader to find.

Wu’s exhaustive search [18] establishes that 1.123732821 is the smallest Perron number of degree 5, so the degree-5 rung is provably optimal among Perron scales. It is also the Lind–Boyd exceptional case: the family $r^d = r + 1$ is Lind’s conjectured smallest Perron number of degree d , corrected by Boyd [17] at exactly $d > 3$ with $d \equiv 3, 5 \pmod{6}$. The exceptional minimisers are finer — 1.1237 against the family’s 1.1673 at $d = 5$ — but they are not one-adder: they cost 3, 5, 7 and more adders as d grows, and beyond $d = 5$ the gain is under 0.8%. A finer rung at five adders is a different trade from a finer rung at one, and only the second would have been news.

Table 31 collects them. Two readings are worth making explicit. Along the one-adder family the adder count is constant and only the registers grow, so the row for r^4 — three adders for a finer ratio than r^5 gives with one — is the exception rather than the pattern, and it is what made an earlier reading of this hierarchy report a cost that grows with degree. And the frequency column does not fall as the ratio does: the $d = 8$ rung measures 815.66 MHz against φ ’s 307.69, because its single adder sits on one of eight parallel paths rather than one of two.

Table 32: A complete ternary neuron, xc7a200t, median of five seeds, harness subtracted, -nodsp. The first column is the implementation this paper first built; the second folds the negation into the accumulation tree, which is where it belonged.

fan-in	LUT/weight	folded	$F_{\max}$	folded	DSP
8	82.5	28.0	109.35	87.29	0
16	87.1	33.1	91.43	73.42	0
32	119.0	33.2	66.53	58.22	0

Theorem 33 (Siegel’s floor bounds the family). *A Pisot–Vijayaraghavan number is a real algebraic integer above one whose conjugates all lie strictly inside the unit circle. Siegel [22] proved the smallest is $1.3247179572\dots$, the root of $r^3 = r + 1$, at any degree and with any integer coefficients. That root is the degree-3 member of the family below, so the family is Pisot at $d = 2$ and $d = 3$ and at no higher degree: every later member is finer than Siegel’s floor and therefore cannot be Pisot.*

Computed conjugate moduli agree exactly: 0.6180 at $d = 2$, 0.8688 at $d = 3$, then 1.0633, 1.0990, 1.0992, 1.0920, 1.0837 at $d = 4\dots 8$. The boundary is not an accident of our enumeration; it is where a theorem from 1944 says it must be, and *the finest scale that is both applicable by one addition and Pisot is the plastic number*.

The distinction matters for what may be built on a rung, not for the arithmetic itself. Normalisation in base θ is realisable by a finite automaton exactly when θ is Pisot [23], and for $n \geq 4$ the root of $x^n = x + 1$ is not even a Parry number, its shift being non-sofic [24]. So a bounded-latency renormaliser or a canonical-form comparator exists at $d = 2$ and $d = 3$ and provably not above. Nothing in this paper’s datapath normalises — the coordinates are carried as integers and never reduced to a canonical digit string — which is why the $d = 8$ step measures 815.66 MHz rather than failing to close. We state it because a reader who assumes a normaliser would be assuming something the mathematics forbids.

Theorem 34 (The one-adder family). *For every d the polynomial $r^d = r + 1$ has a unique real root above one. It has exactly two non-zero coefficients, so its companion map is*

$$(x_0, \dots, x_{d-1}) \mapsto (x_{d-1}, x_0 + x_{d-1}, x_1, \dots, x_{d-2}),$$

one addition regardless of d , and the roots converge to $2^{1/d}$ from above. Any granularity a logarithmic subdivision offers is therefore reachable without a multiplier, at the cost of registers alone.

The map is verified exact over 3000 random coordinate vectors at $d = 5$ and $d = 8$. The convergence is quick: $r/2^{1/d}$ is 1.1441 at $d = 2$, 1.0159 at $d = 5$, 1.0059 at $d = 8$ and 1.0014 at $d = 16$.

The ladder’s cost depends on its role. Applying a ladder weight r^{-j} costs j steps of the companion map. As a *block scale* that is one value per thirty-two weights and the steps amortise; as an *element* it is paid per weight, and a six-bit ladder needs j up to 31. The element form is therefore a coordinate table — the code selects a vector in $\mathbb{Z}[r]$ directly — and Table 34 prices it against what it buys.

The assembled node. Every silicon figure above prices a part. The figure asked from outside is LUTs per neuron, and Table 32 is the assembled thing: fan-in N , weights from the alphabet, accumulator in $\mathbb{Z}[\varphi]$, 8-bit samples, harness subtracted.

Table 33: One pipeline cut in the accumulation tree. The register cost overtakes the frequency gain between sixteen and thirty-two, so the unpipelined form is the better one at fan-in 32 — the opposite of the usual advice, which is why it was measured.

fan-in	LUT/weight	$F_{\max}$ (MHz)	MHz per LUT
8	28.0 $\rightarrow$ 32.4	87.29 $\rightarrow$ 126.73	0.390 $\rightarrow$ 0.489 (+25.6%)
16	33.1 $\rightarrow$ 34.9	73.42 $\rightarrow$ 106.38	0.139 $\rightarrow$ 0.191 (+37.6%)
32	33.2 $\rightarrow$ 47.3	58.22 $\rightarrow$ 75.68	0.055 $\rightarrow$ 0.050 (−8.7%)

Table 34: Element decoders as coordinate tables, isolated, median of five seeds, against the perplexity of the same ladder on SmolLM2-135M (fp32 baseline 14.3607). Thirty-three more LUTs move the error from 70% to 3.7% at no cost in frequency.

decoder	table	LUT	$F_{\max}$	ppl	above fp32
φ , 4-bit	$7 \times 2 \times 5b$	89	621.89	24.4280	+70.1%
$r^5=r^3+1$, 5-bit	$15 \times 5 \times 4b$	95	482.39	15.9242	+10.9%
$r^6=r+1$, 6-bit	$31 \times 6 \times 5b$	122	612.00	14.8882	+ 3.7%

Two’s complement is $(\neg x) + 1$ and the $+1$ of every lane is a single bit, so the fold is an exclusive-or with the weight’s sign plus a narrow second tree carrying the ones, added once at the root. Equivalence was checked before synthesis: 0 mismatches in 400 random (x, w) vectors. Area falls by 2.9–3.6 $\times$ and lands on the hand count — a three-way select of a 16-bit value at roughly 16 LUT6 plus a 16-bit add at roughly 8, so 24 to 30 per lane — which is what confirms the negation was the whole gap rather than part of it. Frequency falls about 20%, the carry tree’s root addition sitting after the main one; on throughput per area the trade is worth 2.35–2.87 $\times$.

So a ternary neuron costs **28.0 LUT per weight at fan-in 8, 33.1 and 33.2 at fan-in 16 and 32, and no DSP at any fan-in.**

And one pipeline register, where it pays. The critical path is the whole adder tree, so cutting it once at the halfway level trades a cycle and $N/2$ registers for frequency — worth measuring rather than assuming, and the answer is not uniform. Equivalence checked first: 0 mismatches in 398 random (x, w) vectors, the pipelined output trailing by exactly one cycle.

At fan-in 8 and 16 the cut buys 45% more frequency for 16% and 5% more area; at 32 it buys 30% for 42%. The crossover is where $N/2$ registers stop being cheap relative to the tree they cut. The configuration to quote is therefore fan-in 8 or 16 with one stage: 32.4 LUT per weight at 126.73 MHz, or 34.9 at 106.38, both at zero DSP. We report the unfolded column beside it because that number was published first, with a caveat naming the negation as its likely cause, and the cause turned out to be the whole of it.

The middle rung is the one a budget would choose: five bits and 95 LUT for +10.9%, a six-LUT premium over φ for a sixfold reduction in error. Applying the scale remains one addition throughout and the datapath remains free of DSP.

Measured as a block scale on two networks, the one-adder ratio matches the geometric ratio beside it at both widths: at five bits 21.5040 against 21.5872 on SmolLM2 and 14.7342 against 14.7814 on Qwen, at six bits 19.3209 against 20.0199 and 13.8528 against 13.8915. The geometric grid’s advantage of Theorem 32 is therefore available with no multiplier at all.

This also resolves the tension left by Theorem 32. The grid minimising relative error is geometric with ratio $R^{1/N}$, an arbitrary real needing a multiplier — in general. Where a block scale actually lives, the occupied range being 8.32 to 9.12 binades and the width four to six

Table 35: One scale application, 16-bit components, xc7a200t, median of five seeds. A fivefold refinement costs 51% more LUTs and no frequency: the eight-register step’s single adder sits on one of eight parallel paths rather than one of two.

step	LUT	FF	$F_{\max}$	level error	adds
$\varphi, r^2 = r + 1$	149	128	307.69	23.607%	1
$r^5 = r^3 + 1$	164	160	306.56	10.575%	1
$r^8 = r + 1$	225	212	815.66	4.625%	1

Table 36: One scale application, xc7a200t, 32-bit, median of five seeds. Degrees 2 and 3 both cost one addition; at degree 4 the coefficient vector stops being sparse and the cost jumps.

scale	LUT	FF	$F_{\max}$	level error	add/sub
φ (deg 2)	223	192	247.10	23.61%	1
plastic (deg 3)	228	200	231.21	13.97%	1
1.1787 (deg 4)	469	320	184.98	8.20%	3

bits, the nearest multiply-free ratio matches the optimum to within 0.06%: 1.551532 against 1.551145 at four bits, 1.236506 against 1.236662 at five, 1.110187 against 1.110180 at six. At eight bits our first enumeration reported a gap, and that was our own search bound rather than a property of the spectrum: stopping at degree 8 the nearest ratio was 1.049852 against a target of 1.026159, while degree 9 reaches 1.036380 and degree 10 reaches 1.027738. Allowing coefficients in $\{0, \pm 1, \pm 2\}$ — still addition and wiring, since ± 2 is a shift — degree 9 reaches 1.026129, within 0.003%, out of 734,533 enumerated roots.

What is a real limit is the one-adder family itself. At the same eight-bit target its best member is 1.0850702 under $\{0, \pm 1\}$ and 1.0510545 under $\{0, \pm 1, \pm 2\}$, errors of 5.6% and 2.4%, and these do not improve with degree. Any ratio is reachable multiply-free if adders are spent; a coarser set is reachable with one.

The hierarchy is cheap only as far as degree three. Table 36 adds the smallest degree-4 scale, the root of $r^4 = -r^3 + r^2 + r + 1$, whose map $(x_0, x_1, x_2, x_3) \mapsto (x_3, x_0 + x_3, x_1 + x_3, x_2 - x_3)$ was verified exact over 3000 random coordinate vectors.

This places the alphabet exactly. φ is the right rung at four bits, where $r^* = 1.4115$ falls between it and supergolden and the two tie. It is the wrong rung at five and above — and by five, the affordable part of the hierarchy has been left behind too, since degree 4 costs $2.1\times$ the area. The regime in which a multiply-free datapath is worth building is the one where degrees 2 and 3 are the only rungs one can pay for, and that regime is three to five bits.

A first predictor failed and the failure is informative: choosing the finest ladder whose span covers the weights’ channel dynamic range — $268.95\times$ median over 155,520 channels — selects the coarsest rung at every budget, correct at three bits and wrong at four and five. Coverage is the wrong criterion because clipping the smallest weights is nearly free while rounding the largest is not. The quantity that decides is total error, and total error is computable before any model runs.

The distinction is between quantizing to *Fibonacci numbers* and working in *the ring those numbers are the coordinates of*. It is narrower than it may sound, and we state it narrowly: by Binet’s formula $F_n = \text{round}(\varphi^n / \sqrt{5})$ exactly for $n \geq 1$ — verified here to $n = 25$ — so a Fibonacci ladder *is* a φ -power ladder up to a constant offset. What does not follow is closure. $F_i F_j$ is not a Fibonacci number, so the product leaves the representable set; $\varphi(a + b\varphi)$ stays in $\mathbb{Z}[\varphi]$. The two schemes agree on the ladder and differ on the ring, and the transformation

Table 37: One accumulation step, xc7a200t, median of five seeds. The non-closed unit is the greedy Zeckendorf normaliser, oracle-checked exhaustively over all 65,536 inputs for the sum, the no-two-adjacent-digits property, and agreement with an independently written greedy model.

one accumulation	LUT	$F_{\max}$ (MHz)
closed, $\mathbb{Z}[\varphi]$, componentwise (16-bit)	182	299.94
non-closed, normalise back (16-bit)	1756	11.46
closed, $\mathbb{Z}[\varphi]$, componentwise (32-bit)	283	239.35
non-closed, normalise back (32-bit)	7017	4.25

stage is the price of the second. $F_i F_j$ escapes the set; $\varphi \cdot (a + b\varphi)$ does not. What the escape costs is exactly the two adders and the routing scheme, and what closure buys is that neither is needed: one integer addition, and the multiply-free path is not an approximation of a multiply but equal to it (Theorem 26).

The closure half of this argument is machine-checked. In the accompanying development the map $(a, b) \mapsto (b, a + b)$ is proved to compute multiplication by φ on the reals, and the exactness of an entire dot product follows from it, in Coq with no admitted lemmas.

The stage costs an order of magnitude, and it runs per accumulation. Corollary 20 says a stage is required; the fabric says what it costs. Table 37 measures one accumulation step in the harness of Section 7 — isolated unit, every output bit folded into the observed reduction, median of five placement seeds.

The stage does not run once per multiply. A product can be handled either way: the identity $F_m F_n = (L_{m+n} - (-1)^n L_{m-n})/5$ turns one into two table lookups and a subtraction. It is *accumulation* that has no such escape, because the sum of two representable numbers is not representable — $F_3 + F_3 = 4 = F_4 + F_1$, and the representation changes non-trivially. At fan-in 512 the stage runs 512 times per neuron, against a pair of integer registers on the closed side.

One limit is worth stating: greedy normalisation is not proved to be the cheapest — constant-time Zeckendorf adders exist — so the ratio is an upper bound on this implementation, while Corollary 20 carries the qualitative claim.

The other apparent limit does not survive measurement. Part of the frequency gap is combinational depth, 46 dependent stages against one addition, and pipelining should convert that into latency and registers. It does, and it is not enough: swept to 32 stages at 16 bits and 48 at 32, frequency rises and then plateaus at 249.44 and 111.36 MHz against the closed path’s 299.94 and 239.35, the second less than half. There is a reason.

Theorem 35 (The normalisation floor). *A normalisation cascade cannot be pipelined below one compare-and-subtract per stage, so its clock period is bounded below by the carry latency of a compare-and-subtract at the datapath width. A compare-and-subtract is strictly more than an addition — a subtraction plus the comparison selecting it — so the floor of a non-closed path lies above the floor of a closed path whose step is a single addition.*

At the plateau the trade is $7.8\times$ and $19.7\times$ the area, $3.6\times$ and $8.3\times$ the registers, 32 and 48 cycles of latency, and still the lower frequency.

This is a different kind of statement from the rest of the paper. Everywhere else we compare error magnitudes between formats; here there is no error to compare. Rounding enters a φ -ternary network only at the nonlinearity, where it is unavoidable in any format, and the cost of exactness in the linear path is two integer accumulators in place of one floating-point one.

Table 38: The zero-DSP stack. One sign-select primitive builds both halves of a node, so the DSP column stays free and radio and network coexist on one part.

block	DSP	result	
GF16 correlator (float taps)	8 DSP48E1	—	before
ternary correlator (GFTernary)	0	78 MHz, real air	after
every transformer block	0	systolic GEMM 33 GOPS	same primitive
trained ternary transformer	0	97% end-to-end	runs
on-silicon backpropagation	0	XOR 4/4, 24–25 epochs	trains

Why the alphabet base cannot be anything else. Closure requires $r^2 = pr + q$ with p, q integers, so that $\mathbb{Z}[r]$ is a ring. Among such $r > 1$ the cheapest recurrence is the one with $p = q = 1$, since any larger p introduces a shift alongside the addition: $r = 1 + \sqrt{2}$ satisfies $r^2 = 2r + 1$ and costs a shift, while $r = \sqrt{2}$ has $r^2 = 2$ and loses the scale out of the lattice entirely. The $p = q = 1$ case is φ , and it is the minimum rather than a member of a set of acceptable choices.

Scope, stated so the claim does not overreach. This concerns arithmetic in a lattice. It applies to the linear algebra that dominates a network’s work and its silicon cost, and it says nothing about control flow, addressing or branching, which are not lattice problems and gain nothing from φ .

23 What the pair costs in silicon

Measured on an XC7A200T through a fully open flow, with no DSP blocks used anywhere:

23.1 The scale applier, and a claim we withdraw

Section 19 closes the multiplier out of the weight and the accumulation. The layer scale is the third site, and our first two attempts at it were each wrong in a different direction. Both are recorded, because the correction is the result.

A ternary network stores weights in $\{-1, 0, +1\}$ and a real per-layer scale $\alpha = \mathbb{E}|W|$. Applying α is a genuine multiply, so the multiplier the ternary weights removed returns at the layer boundary. Snapping α to a grid removes it again, and the question is which grid.

First claim, withdrawn. We reported that the φ^k grid beats powers of two, with excess reconstruction error 2.44% against 4.86% over 210 layers — a ratio of 0.501 against 0.500 predicted from the density $\log 2 / \log \varphi = 1.440$. That measurement reproduces and is correct. It was compared against the wrong baseline. The deployed state of the art for multiplier-free scales is not plain powers of two but **APoT**, whose levels are sums of power-of-two terms. Measured identically, APoT-2 costs 0.1651% and APoT-3 costs 0.0054%, both in one cycle against our k . On accuracy the φ grid loses by $15\times$, and that claim is withdrawn.

Second claim, bounded. Theorem 26 states the linear path is exact because $\mathbb{Z}[\varphi]$ is a ring. It is machine-checked and remains true, but we had presented it as an advantage. An APoT scale is a dyadic rational, and $\mathbb{Z}[1/2]$ is also a ring closed under the datapath’s operations: a fan-in-512 layer with an APoT-2 scale and Q8 inputs accumulates to a denominator of 2^{15} , exact in ordinary fixed point. Binary arithmetic has had this since fixed point existed. The theorem stands; its significance does not.

Table 39: Throughput per area among formats that carry their own range, with conformance. One ternary-neuron datapath, XC7A200T, no DSP, every accumulator bit observed, median of five placement seeds. `int8` is in Table 43. **All three TNF rungs are at their specified parameters here**, and all three moved: the modules previously measured implemented different formats — TNF16 one mantissa bit short, TNF32 with twelve trits and eleven mantissa bits where the specification gives six and twenty-five, TNF64 likewise — and none had a reference, which is how the divergence survived to be quoted. Written from the specification and checked against it they agree on every code swept, and they cost eight to eleven places each. What is left at the top is the alphabet, which agrees on all four of its codes.

	format	storage	LUT	MHz	MHz/LUT	spread	conformance
1	GF Ternary (ours)	2 bits	463	83.19	0.1797	6.61	all 4 codes
2	binary32	32 bits	472	77.00	0.1631	6.37	not swept
3	fp8 e5m2	8 bits	480	67.06	0.1397	5.46	6 of 256 wrong
4	VAX F	32 bits	527	73.51	0.1395	4.43	not swept
5	fp8 e4m3	8 bits	485	67.02	0.1382	6.78	14 of 256 wrong
6	GF 10 (ours)	10 bits	533	72.16	0.1354	5.08	all 1,024
7	binary16	16 bits	522	63.30	0.1213	7.74	all 65,536
8	GF 14 (ours)	14 bits	587	71.15	0.1212	4.54	all 16,384
9	minifloat	8 bits	558	67.06	0.1202	5.49	not swept
10	BNF 16 (ours)	17 bits	571	67.13	0.1176	3.30	all 131,072, 0 wrong
11	TNF 32 (ours)	36 bits	569	66.91	0.1176	6.03	sampled 9,997, 0 wrong
12	TNF 16 (ours)	17 bits	565	66.28	0.1173	8.56	all 131,072
13	TNF 64 (ours)	65 bits	572	65.38	0.1143	8.73	sampled, 0 wrong
14	TNF 8 (ours)	10 bits	545	62.13	0.1140	3.63	all 1,024, 0 wrong
15	GF +8 (ours)	8 bits	628	63.05	0.1004	1.08	selector live; no reference
16	posit8	8 bits	540	44.33	0.0821	4.02	not swept
17	takum16	16 bits	789	58.96	0.0747	5.75	not swept
18	IBM hex32	32 bits	687	46.78	0.0681	3.78	not swept
19	LNS16	16 bits	658	43.04	0.0654	0.45	not swept
20	posit16	16 bits	712	40.38	0.0567	1.02	4 of 65,536 wrong
21	posit32	32 bits	953	28.16	0.0295	2.29	not swept

What measuring area then showed. Both attempts compared accuracy. The appliers differ more in area, and in a way neither of us predicted.

Theorem 36 (Term growth). *A value represented as a sum of n terms from a multiplicatively closed generating set, composed through d scale multiplications without requantisation, has n^{d+1} terms. A value in $\mathbb{Z}[\varphi]$ has two components for every d , since $\varphi^a \varphi^b = \varphi^{a+b}$ adds exponents rather than multiplying terms.*

Theorem 37 (Shifter cost). *A runtime-variable shift over W bits costs $\Theta(W \log W)$ LUTs; a fixed permutation costs zero. An n -term APoT applier is therefore $\Theta(nW \log W)$ and a Fibonacci step is $\Theta(W)$.*

This is the effect we had missed, and it is larger than the one we predicted. We had been calling APoT “shift-add” and treating that as free. A shift by a *constant* is free; a shift by a *variable* is a barrel, measured here at 192 LUTs each on 32 bits, and a scale that must serve any layer is variable by construction. The Fibonacci step contains no shifter at all.

Corollary 22 (Crossover with the multiplier). *Since APoT area grows as n^{d+1} while the multiplier’s does not grow, there is a depth d^* at which APoT costs what it was introduced to avoid. Measured at $W = 32$, $n = 2$: $d^* = 2$, where APoT-8 costs 1217 LUTs against 1215.*

Table 40: Scale appliers, XC7A200T primitives, one harness, DSP inference off. Composition depth d counts scales multiplied without requantisation between them.

applier	LUT	cycles	excess error	
multiplier, real α	1215	1	0.0000%	or 2 DSP48
APoT-2, $d = 0$	384	1	0.1651%	two barrel shifts
APoT-4, $d = 1$	659	1	—	four shifts
APoT-8, $d = 2$	1217	1	—	eight shifts
φ^k , any d	173	k	2.4420%	no shifter

Theorem 38 (No single winner). *In (area, latency, error) no applier dominates another. The multiplier is exact and largest, φ^k smallest and least accurate, APoT-2 between and beaten by neither; all three are vertices of the frontier.*

The recommendation is therefore a rule, not a winner. Use φ^k where the scale path is area-bound, or composed without requantisation — a low-rank factorisation, a folded convolution and batch norm, a residual branch scalar, or accumulation along a mesh route. Use APoT where it is latency- or accuracy-bound. We state this rather than defend one answer because our two previous attempts each defended one, and each was wrong in a different direction.

The area advantage is withdrawn too, and by our own attack. Table 40 was measured in one regime and with one field width, and both choices flattered us.

Theorem 39 (An unrolled recurrence is linear). *A recurrence applier unrolled to depth K over W -bit operands costs $\Theta(KW)$; an additive-term applier whose shifts are compile-time constants costs $\Theta(W)$ independent of the scale, since a constant shift is wiring. Hence when the scale is frozen after training the recurrence loses for every $K \geq 1$.*

Measured: a constant-shift APoT-2 applier costs 26 LUTs, against 64, 128 and 256 for φ^k unrolled at $K = 2, 4, 8$. At the realistic $K = 8$ that is $10\times$ against us, and the lines do not cross even at $K = 1$.

Theorem 40 (Regime inversion). *The area ordering between the two families is a property of neither. It inverts with whether the scale is a compile-time constant or a runtime variable, because a constant shift is wiring and a variable shift is a barrel of $\Theta(W \log W)$.*

Theorem 41 (A barrel is priced by range, not by width). *A runtime shift field need only span the workload’s scale range. Across the 210 layers measured here that range is 3.15 octaves, so two bits suffice. An area comparison that sizes the field larger is measuring a competitor that would never be built.*

At a two-bit field APoT-2 costs 130 LUTs against our 199. The $2.22\times$ advantage reported above was an artefact of a five-bit field, and is withdrawn.

An instrument limitation, stated rather than smoothed. The APoT sweep is non-monotonic — 130, 380, 230, 384 LUTs at shift widths 2 to 5 — with the parameter demonstrably applied and the tool deterministic. This is technology-mapping heuristics, and it means LUT counts from logic synthesis alone are not a reliable area metric at this granularity. A claim resting on a 30% difference between two such points is unsafe; one resting on the $10\times$ of Theorem 39 is.

Table 41: Decode cost, isolated. `nextpnr-xilinx` 1743d0f, XC7A200T, no DSP, median of five seeds. A bare wire in the same harness costs 112 LUT at 827.81 MHz.

	format	kind	LUT	$F_{\max}$
1	GFternary	fixed	66	974.66
2	<code>int8</code>	fixed	76	925.93
3	<code>binary32</code>	fixed	112	886.52
4	TNF16	fixed	101	407.66
6	BNF16	fixed	97	388.35
9	VAX F	fixed	100	311.14
13	<code>binary16</code>	fixed	164	235.18
14	IBM hex32	fixed	243	111.10
15	LNS16	logarithmic	270	93.17
16	<code>posit8</code>	tapered	214	77.75
17	<code>posit16</code>	tapered	302	62.39
18	<code>posit32</code>	tapered	517	49.05

What survives. Not the area advantage. Only Theorem 36 remains: φ^k is the sole applier whose area is independent of composition depth, so at $d = 2$ APoT-8 costs 1217 LUTs where the pair still costs 173. That regime — a chain with no requantisation between links — is real but uncommon. The alphabet’s uniqueness and the $\mathbb{Z}[\varphi]$ closure are machine-checked and untouched, having never been area claims.

The pattern across three attempts at this section is one pattern. Each favourable number came from a comparison whose terms we chose: powers of two rather than APoT, a five-bit field rather than the two the workload needs, a runtime regime rather than the frozen one. The rule that would have caught all three is to write down what the competitor would build if it were trying to win, and measure that.

23.2 The silicon table, measured on an audited harness

Every silicon figure in this paper was originally measured through a harness that combined the decoder with a shared accumulator and observed part of its output. Three separate defects in that arrangement are described in Section 27; each changed a headline number by 10 to 50 percent without any individual measurement being wrong. The table below is measured on the only harness that passes the observation gate we wrote afterwards: decoder alone, every output bit folded into the observed reduction, median of five placement seeds.

Theorem 42 (A taper is paid in delay, not in area). *A tapered format’s decode contains a scan, a serial dependency whose length grows with the word. A serial dependency costs critical path directly and area only incidentally, whereas a radix-heavy fixed format costs a wide shift, which is area-heavy and shallow. The two families therefore separate on frequency and not on area.*

Measured: all fourteen fixed fields sit above all three tapers on frequency, the worst fixed leading the best taper by $1.43\times$; on area they overlap, `posit8` at 214 LUT against IBM hexadecimal’s 243. The overlap is predicted by the theorem rather than excused.

GFternary against `posit32` is $7.8\times$ on area and $19.9\times$ on frequency. GFternary costs 66 LUT where a bare wire costs 112, because decoding a two-bit alphabet lets the synthesiser simplify the downstream register further than passing 32 raw bits does.

The area–frequency correlation on the isolated decoder is -0.60 , against -0.90 on the combined harness. The axes carry independent information here, which is why they are reported separately and no ratio of them appears.

Table 42: Cost of each operation, one harness, XC7A200T, DSP inference off. The logarithmic row is our own measurement of `takum32_decode` from Section 7.

system	multiplication	addition
LNS, takum	exponents add: free	$\log(1+2^x)$ table: 10,967 LUT
fixed point, APoT	multiplier, or wiring if frozen	adder: cheap
$\mathbb{Z}[\varphi]$	Fibonacci step, one adder	componentwise: 64 LUT

23.3 The operating point

The previous subsections withdraw an accuracy claim and an area claim. What remains is a statement about operation profiles, and it is the first in this section that names its own boundary rather than being pushed to it.

A number system is used through two operations, and the three candidates each buy one cheaply and pay for the other.

$\mathbb{Z}[\varphi]$ is the only one cheap under both. Its price is that free multiplication is restricted to powers of φ rather than general, and the following says why that price is not paid in this datapath.

Theorem 43 (The required multiplication set). *In a datapath where every weight is a code applied by sign-select, the layer scale is a power of the alphabet base, and accumulation is the only remaining operation, the multiplications the datapath must perform are exactly $\{\text{base}^k\}$. General multiplication is never requested.*

Corollary 23 (Sufficiency). *A ring closed under addition and under multiplication by powers of its generator is sufficient for such a datapath. $\mathbb{Z}[\varphi]$ is such a ring with both operations at $\Theta(W)$.*

Corollary 24 (The other two are mis-provisioned). *LNS is over-provisioned: it buys general multiplication, which this datapath never asks for, and pays in addition, which it performs on every term of every inner product. Fixed point is under-provisioned in the opposite direction: addition is free, scale multiplication is not, unless the scale is frozen.*

The restriction of $\mathbb{Z}[\varphi]$ coincides with the restriction the ternary datapath already carries. That coincidence is the result; it is not a claim that the format is better, only that it is matched.

Theorem 44 (Compile-time composition is free). *If all d scales of a chain are known before execution, their product is known, and any representation applies it at the cost of a single scale. The term growth of Theorem 36 occurs if and only if d or the scales become known only at runtime.*

This removes three of the four cases where we had claimed a depth advantage. A low-rank factorisation, a folded convolution and batch norm, and a residual branch scalar are all fixed after training, so their product is precomputed and no composition occurs in hardware. One case survives: accumulation along a mesh route without renormalisation, where the hop count is a runtime quantity. The depth advantage is real and it lives there rather than in the network.

Stated boundaries. This is not an area win over APoT, which was measured and withdrawn. It is not an accuracy win, where APoT-2 is $15\times$ ahead. And where a scale is frozen there is nothing to win, because a frozen scale is wiring. It is a claim about which system matches an operation profile, holding where the scale is not frozen: a shared engine serving many layers, or a composition whose depth is not known until execution.

Table 43: Rejected for this task. Each has $e = 0$: no exponent field, hence no dynamic range inside the word. Each functions only alongside an external per-tensor or per-channel scale, and every published sub-8-bit training result using them carries block scales *and* a higher-precision master weight. Gradients span many orders of magnitude; a word with no range cannot hold them.

format	exponent	what supplies the range	trains alone
int8	$e = 0$	learned per-tensor scale	no
<i>measured on the same harness: 454 LUT, 78.83 MHz, 0.1736 MHz/LUT, spread 11.00</i>			
int4	$e = 0$	block scale (32 or 16 elements)	no
int12–int24 accum.	$e = 0$	fixed point, set by hand	no
E2M1 element alone	$e = 2$ in-block	shared UE8M0, outside	no

23.4 The formats a training run cannot live in

Throughput per area is only meaningful among formats that can do the job, so the candidates split in two, by one test that is checkable rather than rhetorical: *does the word carry the range, or must the range be brought from outside?*

Why GFTernary is not in Table 43 despite $e = 0$. The distinction is not the presence of an exponent field but whether the scale is *external and learned* or *intrinsic and exact*. A $\{-1, 0, +1\}$ alphabet needs a real α_ℓ per layer, learned alongside the weights and multiplied by at inference — the multiplier returns. The φ alphabet needs none: by Corollary 19 the inter-layer gain is $F_k\varphi + F_{k-1}$, a pair of integers known in advance, applied by shift and add. Its scale is inside the system by construction rather than fitted from outside it. That is the difference between carrying range and borrowing it, and it is why the stack trains on the part.

Why int8 is reported apart. An **int8** datapath scores 0.180 MHz/LUT at a median of five placement seeds, against GFTernary’s 0.181 — a difference of 0.4% inside a 16% seed spread, so the two are indistinguishable on this flow. An earlier version of this paper reported 0.189 against 0.177 and called it a 7% lead; that came from a single placement run and is withdrawn. **int8** is excluded from that table even so, and the reason is structural rather than favourable to us. **int8** has $e = 0$: no exponent field, hence no dynamic range inside the word. It functions only with an external per-tensor or per-channel scale; without one it is the integers $-128 \dots 127$. By Corollary 27 that is the non-uniformity escape — the scale has been moved outside the word — so **int8** is a block format wearing a scalar’s name, and scoring it against formats that carry their own range measures the wrong thing. It leads the table precisely because it declines the task the table sets.

The consequence matters more for training than for inference. Gradients span many orders of magnitude, and a format with zero dynamic range cannot hold them: every published sub-8-bit training result — MXFP4, NVFP4 — carries block scales *and* a higher-precision master weight. Our own stack trains on the part, at zero DSP, with backpropagation in the format itself.

23.5 Withdrawal: the frontier at equal storage

Corollary 17 and Corollary 26 count positions. A format is stored in bits, and that difference is the whole of the slack.

Theorem 45 (Packed coverage). *A family member realisable in N bits obeys $3^{E_t}2^M \leq 2^{N-1}$.*

Substituting the minimal $E_t = \log_3(b+1)$ gives

$$m + \log_2 3 \cdot \log_3(b+1) = m + \log_2(b+1) \leq N - 1,$$

identically the uniform binary budget of Theorem 47, since $\log_2 3 \cdot \log_3 x = \log_2 x$.

Corollary 25 (The slack is a position artefact). *A ternary exponent encoding provides no budget slack at equal storage. The 0.3691E of Corollary 26 exists only at equal position count, that is on a fabric where a position is physically ternary.*

Measured at equal storage against the catalogue, 6 of 17 sampled formats are dominated rather than all of them. Every uniform binary format spending all its bits ties exactly: binary16 at 15.06 against a budget of 15, binary32 at 31.10 against 31, binary64 at 63.27 against 63. Those dominated are the ones wasting budget — posit32 by 5.12 positions, posit64 by 14.24, takum32 by 0.82, tekum32 by 0.90, `cray_float` by 0.65, and our own `gf8` by 0.70. Several tapered formats escape outright: posit8 by 1.16, posit16 by 2.37, takum16 by 1.53, tekum16 by 1.49. That is consistent with Corollary 27 — escape requires non-uniformity, a taper is non-uniform, and a uniform family cannot catch it at equal storage.

Why this went unseen. The same defect appears twice elsewhere in this paper, as Theorem 8 and as the observation that TNF4 with $E_t = 2$ does not exist in four bits because $3^2 = 9 > 8$. Both were recorded as local facts about packing, and neither was carried back to the result that rests on the same quantity. The rule we adopt is that a correction invalidating a comparison must be applied to every claim resting on that quantity, not only to the one where it surfaced.

23.6 The trit is the slack

Corollary 17 is a fact about a list. It becomes a fact about the design space once one asks what a competitor would have to do to escape.

Theorem 46 (Coverage condition). *$\mathcal{F}_N$ contains a member dominating a format of effective mantissa m and range b binades if and only if*

$$m + \lceil \log_3(b+1) \rceil \leq N - 1.$$

Proof. A member E_t has mantissa $N-1-E_t$ and range $3^{E_t}-1$. Domination requires $N-1-E_t \geq m$ and $3^{E_t}-1 \geq b$, i.e. $\lceil \log_3(b+1) \rceil \leq E_t \leq N-1-m$. Such an integer exists exactly when the stated inequality holds. $\square$

Theorem 47 (Position budget of a uniform format). *Let a uniform format of width N carry a sign, a scale field of E binary positions and a significand of m bits held at every representable magnitude. Then $b+1 \leq 2^E$ and $1+E+m \leq N$, hence*

$$m + \log_2(b+1) \leq N - 1.$$

The two budgets are the same accounting in different currencies, and the exchange rate is the whole result.

Corollary 26 (The ternary exponent is the slack). *Every uniform binary format of width N is dominated by a member of $\mathcal{F}_N$, with slack*

$$\Delta = E(1 - \log_3 2) = 0.3691 E \text{ positions,}$$

strictly positive for $E \geq 1$ and growing linearly in the competitor's own exponent width.

Proof. By Theorem 47, $b + 1 \leq 2^E$ and $m \leq N - 1 - E$. Then $m + \log_3(b + 1) \leq (N - 1 - E) + E \log_3 2 = (N - 1) - E(1 - \log_3 2)$, which satisfies Theorem 46 with the stated margin. $\square$

Measured, in positions: 1.48 at `fp8_e4m3`, 1.85 at `binary16`, 2.95 at `binary32`, 4.06 at `binary64`, 5.54 at `binary128`. The margin is not a coincidence of the catalogue; it is 0.3691 per exponent bit the competitor spends.

This is worth stating plainly because it cuts against Theorem 8, which says a packed ternary exponent never carries more values *per bit* than a binary one, and which we confirmed in silicon: BNF16 and TNF16 differ by 1% in LUTs. Both are true, on different axes. Per *position* a trit spans $\log_2 3$ binades against a bit's one, which is where Δ comes from; per *bit of binary fabric* the packing gives it all back. The ternary encoding earns on the number axis and breaks even on the silicon axis, and a claim that does not name its axis is not a claim.

Corollary 27 (Escape requires non-uniformity). *A width- N format escaping $\mathcal{F}_N$ must violate the hypothesis of Theorem 47: either its significand is not held at every magnitude (a taper), or its scale field is not inside the word (a block).*

The bound is tight, and it is not vacuous. Equality in Theorem 46 is attained by construction, since the width rule spends every position. Among the catalogued formats `posit16` comes within 1%: $10.55 + \log_3 113 = 14.85$ against a budget of 15. And the escape route of Corollary 27 is real rather than hypothetical — the same `posit16`, scored where a taper is designed to live rather than averaged over its range, carries 12 fraction bits while keeping all 112 binades:

$$12 + \log_3 113 = 16.30 > 15.$$

It escapes, as it should: concentrating precision where the workload sits is what tapering is *for*. A bound no format can violate would be measuring nothing.

What this predicts. Corollary 27 names the only two routes out, and both are open questions rather than results of this paper. The taper route is bounded by Theorem 12 and costs what Section 7 measures. The block route — MX and its descendants, where one scale serves 32 elements — lies outside every budget stated here, because the scale is not in the word. We report it as the frontier of this work and not as territory held.

Theorem 48 (Transforms and formats do not compose). *Let $D^*(b, p)$ be the best achievable distortion at b bits per element under a fixed pre-transform p . A transform changes D^* by changing the source it presents; a format chooses a partition attaining some $D \geq D^*$. Two levers that both act on the partition therefore share one ceiling and their gains are sub-additive, while a transform and a format act on different arguments and their composition is order-dependent.*

Measured, the order-dependence reaches 2.78 $\times$: applying the same two levers in the two orders gives different results, which forbids reporting either as a standalone multiplier. We flag this because we made the mistake ourselves — an early version of this work computed a lever's ratio against a losing baseline and reported the levers as nearly multiplicative. They are not.

23.7 Where the block literature is not looking

The block axis is the escape route Corollary 27 leaves open, and it is where the field's attention currently sits. It is worth being precise about what that work varies and what it holds fixed.

The Open Compute microscaling specification fixes a block of 32 elements sharing one UE8M0 scale — a bare power of two, no mantissa — with the element at E2M1. NVIDIA's NVFP4 changes two things: the block shrinks to 16, and the shared scale becomes an E4M3

float rather than a power of two. Both are hardware-resident on current accelerators, which is why the stop condition for this work is stated against them and not against posit or takum.

What the 2026 literature then varies is almost entirely the *transform* applied before quantisation, not the format underneath it: learnable block-wise optimisation for outlier resilience, and rounding fitted to the E2M1 level set, augmented residual channels, and activation-sparsity coupling. Each reports gains, and each holds the element format at E2M1.

That is exactly the variable Theorem 23 addresses. E2M1 spends two of its three non-sign positions on an exponent, *inside a block that already carries a shared scale*. The width rule says the exponent is sized for the range the workload visits, and inside a block normalised by its own maximum that range is measurable rather than assumed. If it is narrower than four binades, E2M1 is buying range the block already paid for, and the positions belong in the mantissa. We measure the within-block span directly and let the rule name a winner before evaluating perplexity; Section 21 reports what came back, including the case where the rule was wrong.

This is a narrower claim than it may look. It says nothing about the transforms, which compose with any element format and whose gains we do not dispute. Our own Theorem 48 says transforms and formats act on different objects — a transform moves the ceiling, a format divides it — so improving the element is orthogonal to improving the rotation, and neither subsumes the other.

24 What this settles

We state the conclusion at the strength the measurements carry, which is a claim about a datapath rather than about a number.

For a ternary datapath — one whose weights are codes and which therefore contains no multiplier — the pair {GFTernary, TNF} is a reference format.

Three things support the word *reference*, and each is a measurement rather than a preference.

It is complete. A ternary node has three format-bearing sites (Section 19); the weight alphabet closes one and the accumulator ladder closes the other, and the sample needs nothing. No other pair in the literature closes both: the ternary-network methods answer the weight and leave the accumulator in a binary float designed for a datapath that multiplies, and the format literature answers the accumulator while assuming a multiplier exists.

It is forced, not chosen. The alphabet base is the unique $r > 1$ with $r^2 = r + 1$ (Theorem 25). The exponent–mantissa split is the solution of a constrained maximisation whose range constraint is active (Theorem 23), so it has no free parameter once the workload’s range is measured. Neither was picked for taste, and both can be re-derived by a reader who disagrees with us.

It is predictive. The error obeys a closed form in the mantissa width alone, so a designer can size the accumulator before building it, and the sizing was checked against real over-the-air capture and against real ternarised weights rather than against synthetic data.

And what it is not. It is not the best format for a datapath that multiplies, where the block axis (MX, NVFP4) holds the ground and we make no claim. It is not a 7% win over `int8`, because `int8` carries no range and is not doing the same job. And the frontier result would be empty without the reason each competitor sits below it, which is why each of those reasons is a theorem here rather than a table entry.

Table 44: Codes lost to packing when a ternary exponent is stored in a binary word. The loss is a property of the integers, not of an implementation.

width	best packed member	codes used	lost
4 bits	$E_t=1, M=1$	6 / 8	25.0%
6 bits	$E_t=3, M=0$	27 / 32	15.6%
8 bits	$E_t=3, M=2$	108 / 128	15.6%
16 bits	$E_t=5, M=7$	31,104 / 32,768	5.1%

25 Where this sits in a sixty-eight-year argument

The claim that three is the better radix is not ours, and stating whose it is makes the contribution of this paper smaller and more defensible.

The people who made the case. Radix economy $\rho(r) = r / \ln r$ is minimised at $r = e$, and three is the nearest integer: ternary sits 0.46% above the optimum where binary sits 6.15% above it. The argument was made mechanically before it was made electronically — Thomas Fowler built a balanced-ternary calculating machine in Devon around 1840, and the recorded motivation was the same economy of digits. It became a working computer in 1958, when **Nikolay Brusentsov** built *Setun* at Moscow State University with the backing of **Sergei Sobolev**, using paired ferrite cores to hold three stable states. Roughly fifty machines were produced between 1959 and 1965. Brusentsov followed it with *Setun-70* in 1970, whose instruction-set ideas anticipated arguments later made independently for RISC; it never reached serial production, and the ternary line was ended administratively rather than technically. **Donald Knuth** kept the idea in circulation, calling balanced ternary “perhaps the prettiest number system” in *The Art of Computer Programming*. The line continues in device work rather than architecture: CNTFET and memristor ternary logic, and more recently silicon-photonic balanced-ternary proposals, all of which build a fabric whose *position* is physically ternary.

What is new here is the boundary, not the claim. We add no support to “ternary beats binary” as a general statement. We measured it three times and it went against us each time — see Theorem 8 and Sections 7 and 23.7. What we contribute is the condition under which the sixty-eight-year-old argument applies, stated exactly:

Theorem 49 (The radix argument holds on positions, not on bits). *Write the economy as cost = $r \times \log_r V$. If the position is the physical unit, the first factor is common to all radices and only the second is compared, giving ternary an unconditional advantage of $E(1 - \log_3 2) = 0.3691E$ positions. If the position must itself be encoded in bits, the format is bounded by $3^{E_t} 2^M \leq 2^{N-1}$, and since 3^{E_t} never divides a power of two the remainder is lost.*

The remainder is arithmetic, not engineering, and it is largest exactly where modern formats live:

Prior work anticipated the mechanism: Parhami’s analysis of binary-encoded balanced ternary treats exactly this encoding penalty. Our contribution is to measure what it costs in a live network and in placed-and-routed silicon rather than in operation counts, and to state it as a condition a designer can check.

Read as a prescription rather than a verdict. The result says what fabric must exist for the radix argument to pay: one where a ternary position is held physically, as *Setun*’s paired cores held it and as CNTFET, memristor and photonic ternary devices propose to hold it again. On such a fabric the 0.3691*E* advantage is collected in full and none of the packing

loss is incurred. On the binary fabric that is actually purchasable today, it is not. That is a narrower claim than the field usually makes, and unlike the broader one it survives our own measurements.

26 What the golden ratio completes

Two lines of argument run through this paper, and they were developed independently and for unrelated reasons. It is worth setting them side by side at the end, because they meet at a single number.

The first line is arithmetic and old. Radix economy makes three the best integer base; Fowler built a machine on it around 1840, Brusentsov and Sobolev built a working computer on it in 1958, and the line has continued into CNTFET, memristor and photonic devices. Section 25 shows why it stalled rather than spread: the moment a ternary position is stored in a binary word, 3^{E_t} fails to divide a power of two and the remainder is lost — a quarter of the alphabet at four bits. The obstruction is arithmetic, so no implementation removes it, and our own three measurements confirm it against us.

The second line is geometric and, in this context, new. If the weights of a network are to be symbols rather than numbers, the datapath must multiply without a multiplier, and that is possible exactly when the product of two symbols falls back into the sum the datapath already forms. Written down, that requirement is $r^2 = r + 1$. It has one positive solution, so the alphabet is not designed but determined: $\{-\varphi, 0, +\varphi\}$ and nothing else (Theorem 25). The consequence is that composition is closed — the gain of k layers is $F_k\varphi + F_{k-1}$, a pair of integers (Corollary 19) — so depth is free of multipliers as well, which the $\{-1, 0, +1\}$ alphabet never achieved because its learned per-layer scale is a real number and must be multiplied by.

The two lines answer each other. What the ternary programme needed was not a better fabric but a base under which its symbols compose; what the golden ratio had, and had had since Euclid divided a line in extreme and mean ratio, was exactly the property $r^2 = r + 1$. Placing it in the weight alphabet turns the sixty-eight-year argument from a claim about storage density into a statement about arithmetic closure — and the measurements in Section 19 follow: zero DSP through a radio, through every transformer block, and through backpropagation on the part itself.

We put this at the end rather than the front because it is a reading of the results and not one of them. Every component is stated and measured above: the uniqueness is two lines of algebra, the Fibonacci closure is arithmetic, the packing loss is a table of integers, and the zero-DSP stack is placed and routed. What they amount to together is left to the reader.

27 Sixteen ways a measurement can be right and its comparison wrong

The retractions in this paper are not incidental to it. Sixteen claims were withdrawn during the work, two of them later withdrawn in turn, and in almost every case the measurement was correct and the comparison around it was not. The regularities that produced them are stated here because they generalise past this subject, and because a reader is entitled to know the shape of the errors a paper has already made.

On what is being compared. A ratio is meaningful only if both sides were built as their own advocate would build them; seven times here a ratio changed sign or magnitude when the other side was rebuilt (Section 23.1). Systems whose range scales differently with width, one exact and one approximate, have no storage-matched ratio at all — a quantity depending on an unstated choice has no limit and cannot converge, which is why six successive instrument

corrections failed to settle one number. Where one system’s figure of merit is constant in a parameter and another’s varies in it, the well-posed comparison is a crossover and not a verdict (Theorem 42 and Section 23.2), and such a crossover is an assertion about the workload rather than about the formats.

On the instrument. Under placement-seed variation LUT count is invariant and achieved frequency moves by 11.4%, so any metric combining them inherits the timing variance in full, and a ranking resolves only rows separated by more than that. A ratio over correlated factors restates the dominant one: on the combined harness, ranking by throughput per area reproduced ranking by $1/\text{area}$ at a Spearman correlation of 0.961. And a result that moves the same way under every successive correction is still moving.

On the harness. Three defects appeared there, none visible in any number, each changing a headline by 10 to 50 percent. A component shared by every candidate is a confound rather than a constant, because it leaves area differences intact and delay differences not. A harness observing part of a design’s output prunes each candidate in proportion to *its own* output width — observing four bits of a 32-bit datapath left 27 LUTs of 179 while leaving an adjacent 6-bit design nearly whole. And net-of-harness subtraction requires the harness to be present in full in every build, since a design consuming fewer of its sources prunes the rest.

On sources. A defect found in an abstract is a hypothesis: abstracts compress, and compression is indistinguishable from omission. An aggregated search result may combine claims from several papers into wording present in none. Both external criticisms we raised during this work dissolved on contact with the papers themselves.

What this is worth. Applied to our own claims the procedure produced sixteen retractions. Applied to two external papers read in full it produced none, and on both occasions it instead caught our own reading of them. We therefore describe it as an instrument for checking one’s own work, and not as a method for auditing anyone else’s. Three of the checks are mechanised as gates in continuous integration; the full statement is in the accompanying note.

27.1 Figures whose generating data is not in the repository

Sixteen claims were withdrawn during this work, and a withdrawn claim’s number does not remove itself from a paper. A traceability check (`tools/check_paper_numbers.py`) extracts every distinctive numeric literal here and asks whether it appears in any file under `research/`, `fpga/` or `conformance/`. It excludes two classes rather than tolerating them: constants stated beside the formula that produces them (60 of them, correctly absent from any data file), and rounding, where a data literal extending the printed digits counts as the source.

Of 451 distinctive literals, 444 trace to a file. **Seven do not**, and they are listed rather than deleted, because a figure that cannot be traced is not thereby shown to be wrong:

The check cannot show a number is right; it finds numbers with no source, which is strictly weaker and still enough. A figure nobody can trace to a file is one nobody can re-check, and after sixteen withdrawals that is exactly where a stale value survives. Two of the seven have a directly measured replacement elsewhere in this paper and disagree with it, which is the pattern this check exists to surface.

Theorem 50 (A document is an artefact, and its pairs need guarding). *Consistency checks are written between artefacts that a tool already reads together — compiler, simulator, build script. Documents are read only by people, so every pair involving one goes unchecked by construction, and that is where withdrawn numbers persist.*

Table 45: Figures with no generating file in the repository. Each is retained, flagged, and should be re-measured before any of them is relied upon.

figure	where it appears
106.68, 215.90, 63.04	extended formats: double-double, quad-double, x87
0.113	takum16 taper slope in the staircase table
0.247	posit32 slope in prose (measured elsewhere here as 0.260)
8.14	one row of the field-width table
8.99	one cell of the by-band effective mantissa table

Enumerating the pairs is what makes a search for such defects systematic rather than accidental. Three earlier gates in this work were each found by tripping over the defect they now catch, and all three guard code against code. Enumerating every pair that must agree exposed five unguarded ones in a single pass, all of them with a document on one side.

27.2 The staircase form prescribes the comparison

Section 4.1 classifies tapers by staircase form to *describe* them. Measuring the crossovers of Section 23.2 over two fitting windows shows the classification does more than that.

Theorem 51 (The staircase form decides whether a window is needed). *A crossover between a fixed field and a taper depends on the fitting window if and only if the taper’s staircase form is not arithmetic. An arithmetic taper’s M_{eff} is linear in $|e|$, so its slope is well-defined without a window. A geometric taper’s is linear in $\log |e|$, so no line fits it and its apparent slope is a property of the window alone.*

Measured across windows 1–9 and 1–20, the split is exact: posit16, posit32 and posit64 move by $1.08\times$, $1.00\times$ and $1.09\times$, while every takum and tekum moves by $1.39\times$ to $1.72\times$.

Corollary 28 (Closed form for a geometric taper). *For a geometric taper, $M(e) = p - \log_2 e$, so the crossover against a fixed field of constant m solves exactly:*

$$x = 2^{p-m},$$

with no window required.

Six figures previously reported as ranges were computed by the wrong model and are corrected here: takum32 moves from 5.9–10.2 to an exact 2.6, tekum16 from 4.2–6.5 to 2.1. The remaining window dependence is confined to posit8, whose peak sits within 0.45 bits of TNF8’s so that any slope estimate divides a very small number.

The point worth keeping is not the arithmetic. **A taxonomy introduced to label formats turns out to prescribe how they must be compared** — arithmetic tapers by a linear crossover, geometric ones by 2^{p-m} . A classification that only labelled would not do this, and we had used it only to label for most of this work.

Theorem 52 (A wobble admits no crossover). *If a format’s M_{eff} is periodic in $|e|$ with amplitude A and period P , and a constant-form competitor’s m lies strictly inside the band, then no x exists such that one leads for all $|e| < x$ and the other for all $|e| > x$. The well-posed statistic is the duty cycle — the fraction of each period on which each leads — together with P .*

Measured on IBM hexadecimal 32, whose M_{eff} by binade reads 21.03, 21.75, 22.94, 19.95, then repeats: amplitude 3.07 bits, period 4 binades, which is $\log_2 16$ as the radix requires. Against packed TNF32 at $M_{\text{eff}} = 21.32$, inside the band $[19.94, 23.11]$, the winner alternates with that period at a duty cycle of 50/50, and no binade separates them.

Table 46: Crossovers, computed by the model each taper’s form requires. A straight-line fit had overstated every geometric taper’s reach, because a line under-reports how fast a logarithm falls near the origin.

taper	form	crossover, binades
posit8	arithmetic	1.5–2.5
posit16	arithmetic	8.5
posit32	arithmetic	9.2
posit64	arithmetic	54.0
takum16	geometric	1.8
takum32	geometric	2.6
takum64	geometric	62.7
tekum16	geometric	2.1
tekum32	geometric	2.4
takum8, tekum8	geometric	TNF everywhere

Corollary 29 (The taxonomy is prescriptive). *Each of the four staircase forms determines the shape of a well-posed comparison against a constant format: a difference for constant, a linear crossover for arithmetic, an exponential crossover 2^{p-m} for geometric, and a duty cycle for wobble.*

The mapping from form to answer-shape is stated here, and we make no claim about how the literature reports. An earlier version of this paragraph said that published comparisons give verdicts or ratios; checked against Hunhold’s takum paper, that is false — it reports by regime with numeric thresholds, noting that posit’s efficiency “markedly diminishes with deviation from unity” and that takum’s accuracy “exceeds that of posits near unity, nearly aligning with it for slightly larger exponents.” The sentence is withdrawn.

What we have not found in that literature is the closed form 2^{p-m} for a geometric taper against a constant-form competitor, or any treatment of the wobble case. The contribution here is the mapping itself: the form of the answer is fixed by the competitor’s staircase, and a classification introduced only to label formats turns out to settle it.

28 Limitations

The fabric this format is optimal on does not exist to buy. Theorem 8 confines the ternary encoding’s benefit to a ternary fabric, and no such process is commercially available. Every hardware figure in this paper is therefore measured on a binary FPGA, where the ternary exponent neither gains nor loses — the enumeration in Section 11 shows it ties a binary encoding exactly at the widths tested. The claims that survive on silicon you can buy are the fixed-field ones: no regime codec, no exponential to evaluate. The ternary claim is architectural and is labelled as such throughout.

1. **The comparison is against takum16, not tekum16.** The oracle labelled tekum decodes all 65,536 16-bit codes identically to the takum oracle. We report the comparison as against takum and note that a genuine tekum comparison remains to be done.
2. **The reference oracle carried a sign defect, now fixed.** Round-tripping 1.5 through the parameterised oracle returned -1.5 at mantissa widths of 21 and 25 bits: the sign bit sat at a fixed position while the exponent mask took more bits than the field holds, so the two overlapped. Both are now derived from the format parameters, and TNF32’s round-trip error falls from $5.4e-1$ — what a systematically inverted sign averages to — to $5.3e-9$.

We record how it survived, because the lesson outlives the bug. The ladder’s check was add/mul commutativity, and an inverted sign survives on both sides of $a + b = b + a$. The check was real and blind to the class. A round-trip assertion sees it, costs one line per probe, and now runs at all nine rungs. tekum [2] is a descendant of takum [1] adapted for balanced ternary. Our oracle reconstructs it from takum’s field scheme; the full per-trit specification requires the tekum paper. The ratios are as good as that model.

3. **No head-to-head in hardware.** takum’s codec RTL is public and written in VHDL [3]. We did not synthesise it alongside TNF, so the cost figures are TNF’s own. What is visible in that source, and which we report as structure rather than measurement: the 16-bit codec instantiates a 725-line generated leading-zero-counter and barrel shifter — the regime decode a fixed-field format does not have.
4. **The range is bounded** at ± 40 in powers of two, roughly ± 12 decades, where a regime field is unbounded. Fixed fields buy the cheap datapath and the uniform mantissa; range is the price, and workloads that need more should raise the trit budget or use a tapered format.
5. **Accuracy above TNF32 is limited by the measuring instrument.** Once the oracle defect was fixed, TNF32 measures $5.3e-9$, flat across the workload, which is what 25 mantissa bits should give. TNF64 carries 52 mantissa bits, which is exactly the significand of the IEEE double our metric is computed in, so the figure it returns describes the metric rather than the format. Measuring the upper rungs requires exact rationals throughout, as the oracle itself already uses.
6. **Five of nine rungs are measured in hardware.** TNF4 through TNF64 fit the device, TNF64 at 7,479 LUTs and 48.20 MHz. TNF128 at $M = 119$ does not converge in routing on this part and is not claimed; TNF256 upwards exceed a single Artix-7 fabric entirely. Table 8 is the specification and the oracle’s coverage; Table 7 is what was placed and routed. We keep them apart on purpose.
7. **One device family.** Artix-7 on an open flow. Not multi-corner characterisation; an ASIC mapping will differ.
8. **The ternary-fabric argument is architectural.** No ternary process exists to synthesise on. Section 7 is a binary-fabric result; the ternary claim is reasoning about regime decode and native exponent addition and should be read as such.

Reproducibility

The reference oracle, the RTL, the equivalence testbenches and the synthesis commands are public. The toolchain is Yosys 0.65, nextpnr-xilinx 1743d0f, Icarus Verilog 13.0 and Python 3.14 — all open-source, none requiring a licence, which is the point: a result nobody can reproduce without buying something is not a public result.

Disclosure

An AI coding assistant was used in preparing the software, the measurements’ tooling and this manuscript. Every figure reported here was produced by running the named tools on hardware in the author’s possession; no measurement in this paper is estimated, and where a figure is derived rather than observed it is labelled.

References

- [1] L. Hunhold. Takum arithmetic. arXiv:2404.18603.
- [2] L. Hunhold. Tekum: balanced-ternary tapered-precision arithmetic. arXiv:2512.10964.
- [3] L. Hunhold. Design and implementation of a takum arithmetic hardware codec in VHDL. arXiv:2408.10594. Source: <https://github.com/takum-arithmetic/Takum-Codec-RTL>.
- [4] D. Vasilev. GoldenFloat: a φ -derived static-split floating-point family from GF4 to GF1024 with a Lucas-exact integer identity. arXiv:2606.05017v3, 2026.
- [5] D. Vasilev. An 83-format numeric catalog with bit-exact conformance vectors: a vendor-neutral reference for FP8, BF16, MXFP4, and microscaling formats. arXiv:2606.09686v2, 2026.
- [6] J. L. Gustafson and I. Yonemoto. Beating floating point at its own game: posit arithmetic. *Supercomputing Frontiers and Innovations*, 4(2), 2017. Posit Standard (2022) fixes $es = 2$ at all precisions.
- [7] C. Wolf. Yosys open synthesis suite. <https://yosyshq.net/yosys/>.
- [8] D. Shah et al. nextpnr: a portable FPGA place-and-route tool. <https://github.com/YosysHQ/nextpnr>; Xilinx target via Project X-Ray.
- [9] B. Hayes, “Third Base,” *American Scientist*, vol. 89, no. 6, pp. 490–494, 2001.
- [10] Radix economy, optimal radix choice. https://en.wikipedia.org/wiki/Radix_economy
- [11] Ternary27: A Balanced Ternary Floating Point Format, Version 3.1. <https://hackaday.io/project/183153>
- [12] D. Goldberg, “What Every Computer Scientist Should Know About Floating-Point Arithmetic,” *ACM Computing Surveys*, vol. 23, no. 1, pp. 5–48, 1991.
- [13] IBM hexadecimal floating-point. https://en.wikipedia.org/wiki/IBM_hexadecimal_floating-point
- [14] P. Elias, “Universal codeword sets and representations of the integers,” *IEEE Trans. Inform. Theory*, vol. 21, no. 2, pp. 194–203, 1975.
- [15] Closing the Gap Between Float and Posit Hardware Efficiency, arXiv:2603.01615.
- [16] Performance-Efficiency Trade-off of Low-Precision Numerical Formats in Deep Neural Networks, arXiv:1903.10584.
- [17] D. W. Boyd, “The maximal modulus of an algebraic integer,” *Math. Comp.*, vol. 45, no. 171, pp. 243–249, 1985.
- [18] Q. Wu, “The smallest Perron numbers,” *Math. Comp.*, vol. 79, no. 272, pp. 2387–2394, 2010.
- [19] C. Frougny, E. Pelantová, and M. Svobodová, “Minimal digit sets for parallel addition in non-standard numeration systems,” *J. Integer Seq.*, vol. 16, art. 13.2.17, 2013.
- [20] J. Légerský, “Construction of algorithms for parallel addition in expansions in an algebraic base,” *Acta Polytechnica*, vol. 58, no. 5, pp. 285–291, 2018.

- [21] C. Frougny, E. Pelantová, and M. Svobodová, “Parallel addition in non-standard numeration systems,” *Theoretical Computer Science*, vol. 412, no. 41, pp. 5714–5727, 2011.
- [22] C. L. Siegel, “Algebraic integers whose conjugates lie in the unit circle,” *Duke Math. J.*, vol. 11, no. 3, pp. 597–602, 1944.
- [23] D. Berend and C. Frougny, “Computability by finite automata and Pisot bases,” *Math. Systems Theory*, vol. 27, pp. 275–282, 1994.
- [24] S. Akiyama, “A family of non-sofic beta expansions,” *Ergodic Theory Dynam. Systems*, vol. 36, no. 2, pp. 343–354, 2016.
- [25] V. Egiazarian et al., “Bridging the Gap Between Promise and Performance for Microscaling FP4 Quantization,” arXiv:2509.23202, 2025.
- [26] NVIDIA, “Introducing NVFP4 for Efficient and Accurate Low-Precision Inference,” 2025; and “Pretraining Large Language Models with NVFP4,” arXiv:2509.25149, 2025.
- [27] J. Lin, J. Tang, H. Tang, S. Yang, W.-M. Chen, W.-C. Wang, G. Xiao, X. Dang, C. Gan, and S. Han, “AWQ: Activation-aware Weight Quantization for On-Device LLM Compression and Acceleration,” in *Proc. MLSys*, 2024. arXiv:2306.00978.
- [28] X. Yang, Y. Wang, Y. Qin, J. Wang, S. Wei, Y. Hu, and S. Yin, “FQP: A Fibonacci Quantization Processor with Multiplication-Free Computing and Topological-Order Routing,” in *Proc. 61st ACM/IEEE Design Automation Conference (DAC)*, 2024. <https://doi.org/10.1145/3649329.3656502>
- [29] R. Fiandaca and M. D. Gomony, “Fibbinary-Based Compression and Quantization for Efficient Neural Radio Receivers,” arXiv:2511.01921, 2025.
- [30] S. Ma et al., “The Era of 1-bit LLMs: All Large Language Models are in 1.58 Bits,” arXiv:2402.17764, 2024.

Figure 6: $\kappa(r)$ against the scale radix. The unconstrained minimiser sits near $r = 1.9$ and is unreachable: only $r = 2^k$ scales by a shift, and on that set κ is strictly increasing.




Figure 7: The same nineteen formats. Ours are the four families named above; the bar labels carry the LUT count so area and throughput are readable together.